

# SQL 1: CREATE and DROP Tables

Databases and Interfaces

---

Matthew Pike & Yuan Yao

University of Nottingham Ningbo China (UNNC)

## Overview

---

In this lecture, we will look at:

- Review - what is a DBMS?
- What is SQL? What does it allow us to do?
- How we can use SQL to:
  - **CREATE** tables in a database
  - Link tables together using **FOREIGN KEY** constraints
  - **DROP** (delete) a table from a database

## DBMSs and SQL

---

# Database Management Systems (DBMS)

A DBMS is a collection of programs that enables users to create, maintain, and interact with a database. Some key aspects of a DBMS include:

- A structured way to organise, store, and retrieve data.
- A language (often - SQL) to query and manipulate the data in the database.
- An administrative interface (often a CLI) to manage the DBMS.
- A programmatic interface (API) for applications to interact with the database.
- Critical functions like security, concurrency control, transaction management, crash recovery - to preserve data integrity.
- Examples of DBMSs, include:
  - SQLite
  - MariaDB
  - MySQL

# SQL - Structured Query Language

SQL is a standard language for managing data in a relational database which builds upon Edgar F. Codd's relational model (Codd 1970). Some key aspects of SQL include:

- SQL is a declarative language, meaning that you specify what you want, not how to get it.
  - This is in contrast to imperative languages, which require you to specify the exact steps to achieve your desired outcome.
- Statements are not necessarily run/executed in the order they are written.
  - There are however rules about the order in which statements are declared.
- Example:
  - `SELECT * FROM Student;` - Retrieve all data from the **Student** table
  - `SELECT * FROM Student WHERE SID < 100;` - Retrieve all data from the **Student** table where the **SID** is less than 100.

- SQL became a standard of the:
  - American National Standards Institute (ANSI) in 1986;
  - International Organization for Standardization (ISO) in 1987.

## **i** SQL and DBMS Interoperability

Keep in mind that while SQL is a standard, it is not supported *exactly* the same way by all DBMSs. In practice, you may need to update your SQL queries to work with different DBMSs.

Consider:

*“The folding of unquoted names to lower case in PostgreSQL is incompatible with the SQL standard, which says that unquoted names should be folded to upper case. Thus, **Foo** should be equivalent to **F00** not **foo** according to the standard.” (Wikipedia 2023)*

SQL consists of many types of operations for creating, selecting, updating and removing data in the database. Informally, we can divide SQL into three sublanguages:

1. **Data Definition Language (DDL)** - used for creating and modifying database objects, such as tables, indices, and other structural elements. DDL statements define the structure and organisation of the data in the database.
2. **Data Manipulation Language (DML)** - DML is used for inserting, retrieving, and manipulating data in a database.
3. **Data Control Language (DCL)** - DCL is used for controlling security and concurrent access to a database. It includes statements for granting and revoking privileges, creating and dropping user accounts, and managing transaction locking and isolation.



## Creating Tables with CREATE in SQL

---

# Terminology

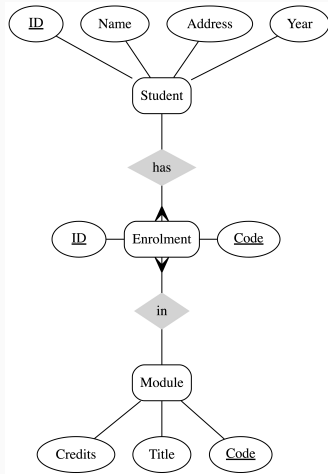
- We have already looked at relational and ER representations of data.
- Now, we will look at how to realise these designs in a real (relational) database, using SQL.
- Table 1 provides a mapping of the terminology used between different representations.

Relations	ER Diagrams (ERD)	Relational Databases
Relation	Entity	Table
Tuple	Instance	Row
Attribute	Attribute	Column/Field
Foreign Key	M:1 Relationship	Foreign Key
Primary Key	<u>Attribute</u>	Primary Key

**Table 1:** Terminology mapping between Relational, ERD and Relational Databases

# Going From ERD to Relational Databases using SQL

- Goal:
  - Given an ERD (such as Figure 1), create a relational database using SQL to represent the structure of the data.
- To do this, we need to:
  1. Translate Entities into Tables.
  2. Translate Attributes into Columns.
  3. Approximate attribute domains by assigning data types to Columns.
  4. Translate relationships into Foreign Keys.



**Figure 1:** ERD for Student Module Enrolment

## Example: Student Table

### Goal

Create a table in SQL to represent the **Student** entity in Figure 2. Student IDs are unique and cannot be **NULL**. Addresses are optional and can be **NULL**. If not specified, the **Year** of study defaults to 1.




Figure 2: ER Diagram for Student Table

## Step 1: Translate Entities to Tables

```
CREATE TABLE Student(  
    ...  
);
```

## Step 2: Attributes of an Entity become Columns

```
CREATE TABLE Student (  
    sID ,  
    sName,  
    sAddress,  
    sYear  
);
```

## Step 3: Assign Types to Columns

```
CREATE TABLE Student (  
    sID INTEGER,  
    sName VARCHAR(50),      -- Reasonable?  
    sAddress VARCHAR(255),  -- Reasonable?  
    sYear INTEGER  
);
```

### Comments in SQL

Just as with other programming languages, SQL supports comments. Comments are ignored by the DBMS and are used to document your code.

- Single line comments start with --
- Multi-line comments start with /\* and end with \*/

## Step 4: Add constraints

### Note

Both SQL statements below are equivalent to one another.

```
CREATE TABLE Student (  
    sID INTEGER PRIMARY KEY,  
    sName VARCHAR(50) NOT NULL,  
    sAddress VARCHAR(255),  
    sYear INTEGER DEFAULT 1  
);
```

```
CREATE TABLE Student (  
    sID INTEGER,  
    sName VARCHAR(50) NOT NULL,  
    sAddress VARCHAR(255),  
    sYear INTEGER DEFAULT 1,  
    CONSTRAINT pk_student PRIMARY KEY (sID)  
);
```



# Constraints

- Constraints are an essential aspect of database design, as they enforce rules on the data stored in a table.
  - These rules ensure that the data is consistent, accurate, and reliable.
- For example:
  - Constraints can be used to specify that a column cannot contain **NULL** values, or that all values must be **UNIQUE**.
- You can specify a name for each constraint, which makes it easier to reference and manage them.
  - We saw in the previous example - a constraint named **pk\_student**.
  - If you don't specify a name, one will be generated for you.
  - Naming constraints is good practice and should be done whenever possible.

## Primary Key and Unique Constraints

In SQL, **PRIMARY KEY** and **UNIQUE** constraints are used to enforce uniqueness and non-nullness on columns or sets of columns in a table.

- A **PRIMARY KEY** constraint uniquely identifies each row in a table. It is a column or set of columns that cannot contain **NULL** values and must contain unique values for each row.
- A **UNIQUE** constraint ensures that all values in a column are different.
  - It is similar to a **PRIMARY KEY** constraint, but it can contain **NULL** values.



SQLite allows **NULL** values in **PRIMARY KEY** columns!

“According to the SQL standard, **PRIMARY KEY** should always imply **NOT NULL**. Unfortunately, due to a bug in some early versions, this is not the case in SQLite. Unless the column is an **INTEGER PRIMARY KEY** or the table is a **WITHOUT ROWID** table or a **STRICT** table or the column is declared **NOT NULL**, SQLite allows **NULL** values in a **PRIMARY KEY** column.”

- From [https://www.sqlite.org/lang\\_createtable.html#primkeyconst](https://www.sqlite.org/lang_createtable.html#primkeyconst).

## Types in SQL

---

## Data Types are DBMS Dependent

Not all data types are supported by all DBMSs, and some data types may be implemented differently by different DBMSs.

- SQL provides a number of data types for representing data in a database.
- These include:
  - Numeric types: **INTEGER, REAL, NUMERIC**
  - Character types: **CHAR, VARCHAR(M)**
  - String types: **VARCHAR, TEXT**
  - Date and time types: **DATE, TIME, TIMESTAMP**

## Examples of Data Types

Data Type	Description	Example
INTEGER	Integer value	1, 2, 3
REAL	Floating point value	1.0, 2.0, 3.0
CHAR	Fixed length string	'a', 'b', 'c'
VARCHAR or TEXT	Variable length string	'a', 'ab', 'abc'
DATE	Date value	'2018-10-01'

**Table 2:** Examples of data types in SQL

## SQLite Types

More information on SQLite types can be found: <https://www.sqlite.org/datatype3.html>

- Most SQL DBMSs uses *static*, rigid typing.
  - With static typing a value's datatype is determined by the column in which the value is stored.
- SQLite uses a more general *dynamic* type system.
  - The datatype of a value is associated with the value itself, not with its column's datatype.
- SQLite 3 defines 5 affinity types, to which a column's datatype will be assigned:
  - TEXT, NUMERIC, INTEGER, REAL, BLOB

## Another Example: Module Table (1/2)

### **i** Module Table

The **Module** table stores information about modules offered by the university. Each module has a unique 8 character module code, a title and a credit value.

```
CREATE TABLE Module (  
    ...  
);
```




Figure 3: ER Diagram for the Module Table

## Another Example: Module Table (2/2)

### The **DEFAULT** clause

The **DEFAULT** clause can be used to specify a default value for a column. If no value is specified for a column when a new row is inserted into the table, the default value will be used instead.

```
CREATE TABLE Module (  
    mCode CHAR(8) PRIMARY KEY,  
    mTitle VARCHAR(100) NOT NULL,  
    mCredits INTEGER NOT NULL DEFAULT 10  
);
```

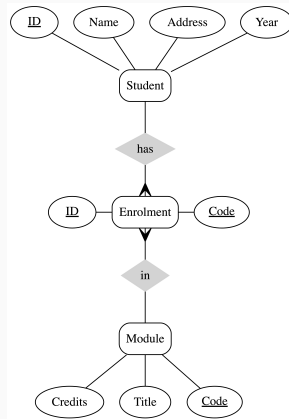


## Relationships

---

## Example: Student-Module-Enrolment

- Currently, we have two tables:
  - **Student**
  - **Module**
- We need to add a table, **Enrolment** to represent the relationship between **Student** and **Module**.
- This table will have two columns:
  - **sID** – references the primary key in the **Student** table.
  - **mCode** – references the primary key in the **Module** table.



**Figure 4:** ERD for the Student, Module and Enrolment example

# Foreign Keys

- Foreign keys are used to create *relationships* between tables.
- **M:1** relationship: Represented by a foreign key in the *many* table.
- **M:M** relationship: are split into two **1:M** relationships.
  - A table is used to represent the relationship between the two tables.
  - This table is called a *link* or *junction* table.
- Why Foreign Keys are important:
  - Relationship building - Allow us to link data between tables.
  - Data Integrity - Relationships between tables are accurate and consistent.
  - Data Consistency - Data updates are propagated to all related tables.



Foreign Keys must reference a **UNIQUE** (typically **PRIMARY KEY**) column

A foreign key must reference a **UNIQUE** column in the referenced table, otherwise foreign key constraints cannot be enforced.

## Example: Add Columns to Enrolment Table

```
CREATE TABLE Enrolment (  
    sID INTEGER NOT NULL,  
    mCode CHAR(8) NOT NULL  
    ...  
);
```

### Module Table Definition


We haven't defined the **Module** table yet in the lecture slides. We will do this later in the lecture.

## Example: Adding Foreign Keys

```
CREATE TABLE Enrolment (  
    sID INTEGER NOT NULL,  
    mCode CHAR(8) NOT NULL,  
    -- Composite Primary Key  
    PRIMARY KEY (sID, mCode),  
    -- Specify that sID is a foreign key  
    FOREIGN KEY (sID)  
        -- References the Student table  
        REFERENCES Student(sID),  
    FOREIGN KEY (mCode)  
        REFERENCES Module(mCode)  
);
```

- The **FOREIGN KEY** constraint specifies that the values in the column(s) must match values in the referenced column(s).
- The **REFERENCES** keyword specifies the table and column(s) that the foreign key references.
- The referenced column(s) must be a **PRIMARY KEY** or have the **UNIQUE** constraint.

## Visualising Foreign Key Relationships



**Figure 5:** Visualisation of the foreign key relationships between the *Student*, *Module* and *Enrolment* tables.



## SQLite Foreign Key Constraints

By default, SQLite does not enforce foreign key constraints. You need to enable them using the **PRAGMA** statement:

```
PRAGMA foreign_keys = ON;
```

- Referential integrity constraints can be specified for each foreign key
- When relations are updated or deleted, constraints are checked
- There are three options:
  - **RESTRICT**: The database will not allow the update or delete to proceed if it would break referential integrity
  - **CASCADE**: The database will update/delete related rows in the other table
  - **SET NULL**: The database will set the foreign key to **NULL** in the related row in the other table

## Example: Add Referential Integrity Constraints

```
CREATE TABLE Enrolment (  
    sID INTEGER NOT NULL,  
    mCode CHAR(8) NOT NULL,  
    PRIMARY KEY (sID, mCode),  
    CONSTRAINT en_fk1  
        FOREIGN KEY (sID) REFERENCES Student(sID)  
        ON UPDATE CASCADE  
        ON DELETE CASCADE,  
    CONSTRAINT en_fk2  
        FOREIGN KEY (mCode) REFERENCES Module(mCode)  
        ON UPDATE CASCADE  
        ON DELETE CASCADE  
);
```



## Deleting Tables using DROP

---



### Practice Caution using **DROP**

Be **very careful** with this command. It will delete the table and all data. There is no undo.

- You can delete tables with the **DROP** keyword:
  - `DROP TABLE [IF EXISTS] table-name;`
- For example:
  - `DROP TABLE IF EXISTS Student;`
- Foreign Key constraints will prevent you from deleting a table if it is referenced by another table.
  - You can delete the referencing table first, then the referenced table
  - Although, by default, SQLite does not enforce foreign key constraints. You need to enable them using the **PRAGMA** statement:
    - `PRAGMA foreign_keys = ON;`

## Reference Section

---

## CREATE Table Definition

```
CREATE TABLE table-name (  
    col-name-1 col-def-1,  
    col-name-2 col-def-2,  
    ...  
    col-name-n col-def-n,  
    constraint-1,  
    ...  
    constraint-k  
);
```

- `table-name` is the name of the table to be created
- `col-name-n` is the name of the n-th column
- `col-def-n` is the definition of the n-th column
- `constraint-k` is the k-th constraint on the table

# CREATE Column Definition



## Non-Exhaustive List of Column Constraints

More information: [https://www.sqlite.org/lang\\_createtable.html](https://www.sqlite.org/lang_createtable.html)

```
col-name col-def  
[NULL | NOT NULL]  
[DEFAULT default_value]  
[NOT NULL | NULL]  
[AUTO_INCREMENT]  
[UNIQUE]  
[PRIMARY KEY]
```

- `col-name` is the name of the column
- `col-def` is the definition of the column
- `NULL` or `NOT NULL`: whether the column can contain `NULL` values
- `DEFAULT default_value`: specifies a default value for the column
- `AUTO_INCREMENT`: column is an auto-incrementing integer
- `UNIQUE`: must contain unique values
- `PRIMARY KEY`: column is a primary key

## Foreign Key Constraints

```
CONSTRAINT name
  FOREIGN KEY
    (col1, col2, ...)
  REFERENCES
    table-name
    (col1, col2, ...)
ON UPDATE ref_opt
ON DELETE ref_opt
```

- You need to provide:
  - A name for the constraint
  - The name of the column(s) in the referencing table
  - The name of the table being referenced
  - The name of the column(s) in the referenced table
  - The action to take when the referenced row is updated
  - The action to take when the referenced row is deleted
- `ref_opt` can be : `RESTRICT` | `CASCADE` | `SET NULL` | `SET DEFAULT`

### SQLite dot commands

More information: <https://www.sqlite.org/cli.html>

- The SQLite Command Line Interface (CLI) has special commands dot commands `.`
- `.` commands control the behaviour of the CLI
- The most useful commands are:
  - `.help` - Display a list of commands
  - `.tables` - Display a list of tables
  - `.import` - Import data from a file into a table
  - `.read` - Execute commands from a file
  - `.schema` - Display the schema of a table
  - `.quit` - Exit the command line tool

## Extra-Study Exercise: Pilot Qualification Database

### 💡 Problem Description

A pilot can be qualified to fly multiple aircraft, and an aircraft can be flown by many pilots. All pilots must have a name and age. All pilots begin with 1 year of experience (from training). All aircraft must have all attributes.

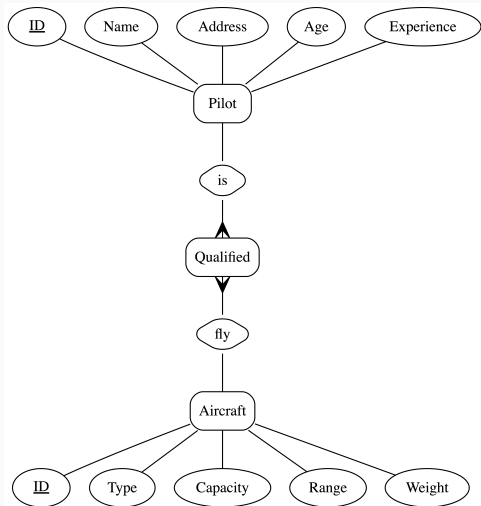


Figure 6: ERD for the Pilot Qualification example



Codd, Edgar F. 1970. "A Relational Model of Data for Large Shared Data Banks." *Communications of the ACM* 13 (6): 377–87.

Wikipedia. 2023. "SQL — Wikipedia, the Free Encyclopedia." <https://en.wikipedia.org/wiki/SQL>.