



University of
Nottingham

UK | CHINA | MALAYSIA

Operating Systems and Concurrency

Lecture 1: Introduction

University of Nottingham, Ningbo China
2024



Module Convenors

- Week 2 – Week 9
 - Name: Dr. Fiseha Berhanu Tesema
 - Office: PMB-442 (by appointment)
 - Office Hour: Moodle Page
 - E-mail: Fiseha-Berhanu.Tesema@nottingham.edu.cn
- Week 10 – Week 14
 - Name: Dr Qian Zhang
 - Office: SEB-438, Office hour (physical)
 - Office Hour: Moodle Page
 - E-mail: qian.zhang@nottingham.edu.cn

Goals

What and How?

- **Goals:**

- Introduction to the **fundamental concepts, key principles** and **internals** of operating systems and concurrency.
- Better understand how **application programs interact/rely with the operating system**
- Basic understanding of writing **concurrent/parallel code**.

- **How:**

- Through **lectures** (2hrs/w for 11 weeks, 3hrs/w for one week(Week 14th))
 - Monday, DB-C05, 15:00 – 17:00 (2 hour session)
 - Tuesday, DB-C05, 13:00-14:00 (1 hour session)
- Through **labs/coursework** (lab starts from 25th September)
 - Wednesday, IAMET-406, 11:00-13:00 (2 hour session)
 - Students from both sessions arrive at 11:00, take attendance according to your timetable
 - There will be in-lab quizzes during the 1st, 6th, and 12th week. These quizzes are worth 0%, 2%, and 3% respectively, contributing to your total mark.



Goals

What and How?

- **Lectures** will **introduce** the concepts
- The **labs** will teach you:
 - Practical experience/insights in fundamental **OS concepts**
 - The use of operating system APIs and implementation of **OS schedulers**
 - The basics of **concurrency**
 - Help you with practical **implementations/coding** (e.g., concurrent programs)

- A **120 minute exam** that focusses on:
 - Knowledge
 - Comprehension
 - Application
- The exam will be 3 out of 4 questions, with **75%** of the assessment on the exam.
- **Sample questions** from previous years are available on Moodle and will be included in the lectures
- The **coursework** is worth **20%**
 - Coursework uses the **concepts** introduced in the labs
- **In-Lab Quiz** is worth **5%**



Content

Subjects We Will Discuss

Subject	# Lectures
Introduction to operating systems/computer design	2
Processes, process scheduling, threading, ...	4-5
Concurrency, deadlock	4-6,2
Memory management, swapping, virtual memory,...	6-7
File System, file structure, management,...	4-5
Revision	1

Table: Preliminary course structure


Reading Material

- Books:

- Tanenbaum, Andrew S. 2014 **Modern Operating Systems**. 4th ed. Prentice Hall Press, Upper Saddle River, NJ, USA.
- Silberschatz, Abraham, Peter Baer Galvin, Greg Gagne. 2008. Operating System Concepts. 8th ed. Wiley Publishing.
- Stallings, William. 2008. **Operating Systems: Internals and Design Principles**. 8th ed. Prentice Hall Press, Upper Saddle River, NJ, USA.
- Thomas Anderson and Michael Dahlin. 2014 Operating Systems, Principles & Practice. 2nd Ed. Recursive Books, Ltd.

- Other sources:

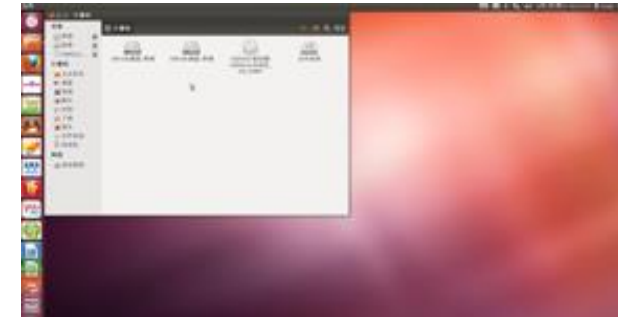

- Daniel P. Bovet, Marco Cesati Understanding the Linux Kernel. 3rd ed. O'Reilly Media, November 2005
- **Course slides** will be available on Moodle



- **Defining** operating system
- What is **multi-programming**
- **Kernel-user mode** and **system calls**

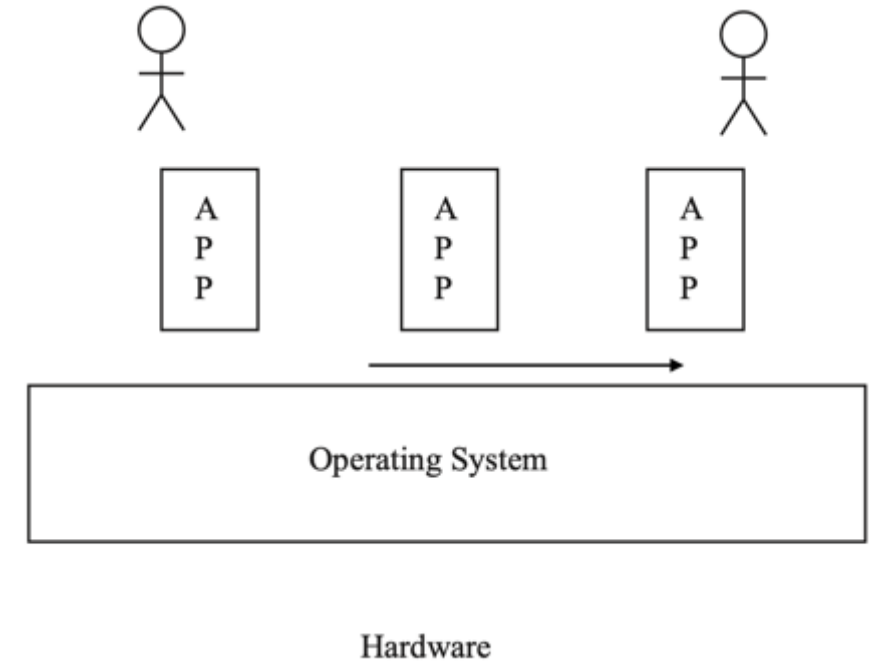
What is Operating System

- Modern computer **consists of processor, main memory, disks, printers, keyboard, mouse, display and various other input/output (I/O) devices.**
- Operating system is a layer of software to provide user programs with a **better, simpler, cleaner** model of the computer and to handle **managing all the resources.**
- User normally interact with it **using user interface** programs like Shell and GUI (Graphical User Interface)



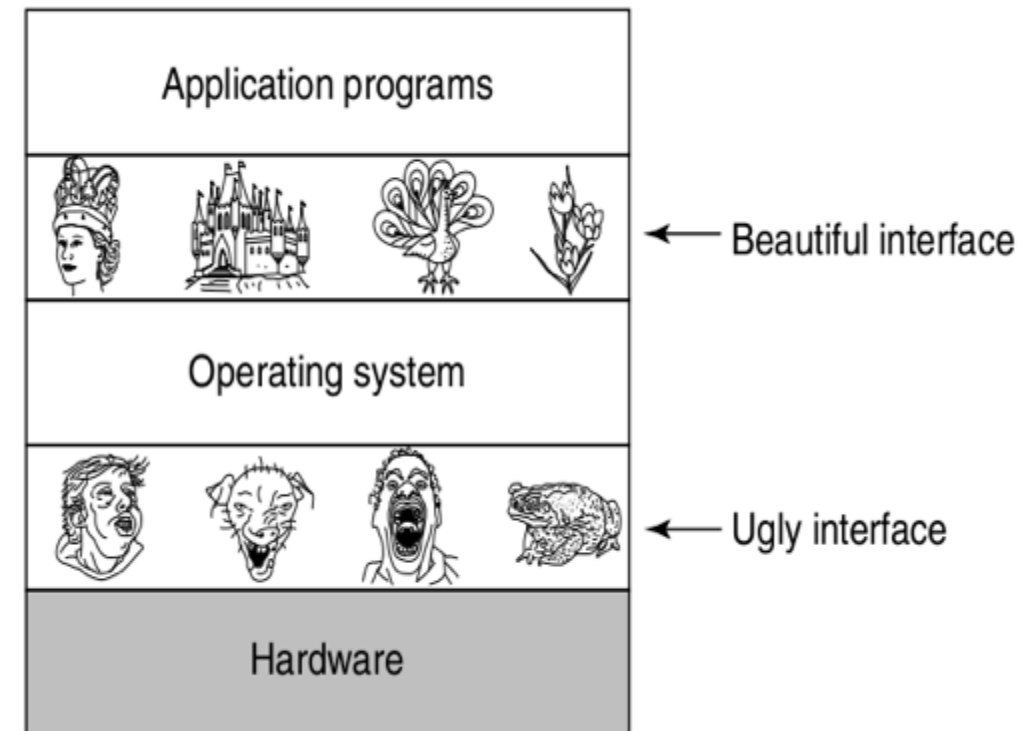
What is Operating System

- For Application Programmers
 - It is an **extended machine**
 - Hides the messy details which must be performed
 - Presents user with a **virtual machine**, easier to use
- For OS Kernel Developers
 - It is a **resource manager**
 - Each program **gets time** with the resource
 - Each program **gets space** of the resource



Defining Operating Systems

- An operating system is a layer of **indirection** on top of the hardware:
 - It provide **abstractions** for application programs (e.g., **file systems, process, address space**)
 - It provides a **cleaner** and **easier interface** to the hardware and hides the complexity of 'bare metal'
 - It allows the programmer to be lazy by using common routines





What is Operating System


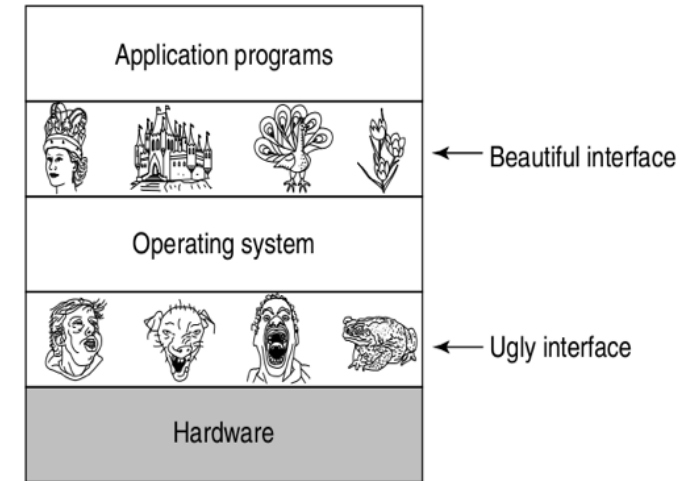
What can an OS do for me?

- Where is the file physically written on the disk and how is it retrieved (**file systems**)
- Why do instructions look the same regardless of the device? (**abstraction**)
- What if multiple programs access the same file simultaneously? (**processes, concurrency,...**)
- Why is the access denied? (**protection**)

Defining Operating Systems

A virtual Machine Providing Abstractions

- **Challenges of Early Programming:**
 - Programmers had to **manually** control each piece of hardware, such as memory, input/output devices, and the CPU, often needing to know the specific characteristics of the hardware they were working with.
 - **(Ugly Hardware)** Real computer hardware is intricate, inconsistent, and difficult to control. It requires precise timing, memory management, and handling of peripheral devices like keyboards, printers, and displays.





Defining Operating Systems

A virtual Machine Providing Abstractions

- **Challenges of Early Programming:**
 - Adding Two Numbers in Assembly (x86 architecture)

```
MOV AX, [1000h] ; Load the value from memory address 1000h into register AX
MOV BX, [1001h] ; Load the value from memory address 1001h into register BX
ADD AX, BX      ; Add the values in AX and BX, store the result in AX
MOV [1002h], AX ; Store the result from AX into memory address 1002h
```

How This Relates to Early Programming:

- The programmer is directly **manipulating CPU registers and working with specific memory addresses**.
- They need to understand **how the memory is laid out and how the CPU handles data in registers**.
- The programmer is **manually** loading values from specific memory addresses, performing operations on them, and then storing the result in another memory address.
- **Error-Prone:** If the programmer accidentally used the wrong memory address, or made a mistake in the ADD or MOV instructions, the program could fail, or worse, overwrite critical parts of memory.


Defining Operating Systems

A virtual Machine Providing Abstractions

- A good abstraction serves two primary functions:

1. Define and Implement Abstractions:


- This involves creating **clear, standardized models of complex hardware or software systems** that hide the underlying details. e.g.:
 - **File Systems:** Rather than dealing with bits on a disk, an operating system **abstracts this into files and directories**.
 - **Processes:** Instead of directly **managing the CPU, memory, and devices** for each program, the OS abstracts these resources into manageable **units called "processes."**
 - **Memory Management:** **Virtual memory** abstracts the complex physical memory layout, allowing programs to think they have access to **a continuous block of memory**.



Defining Operating Systems

A virtual Machine Providing Abstractions

- A good abstraction serves two primary functions:
 2. Use Abstractions to Solve Problems:
 - Once the abstractions are defined, programmers can use them to focus on solving their specific problems **without worrying about the hardware complexity.**
 - This allows developers to:
 - Write software more quickly and efficiently.
 - Build complex systems with reusable components.
 - Leverage system resources in a more structured and manageable way.





Defining Operating Systems

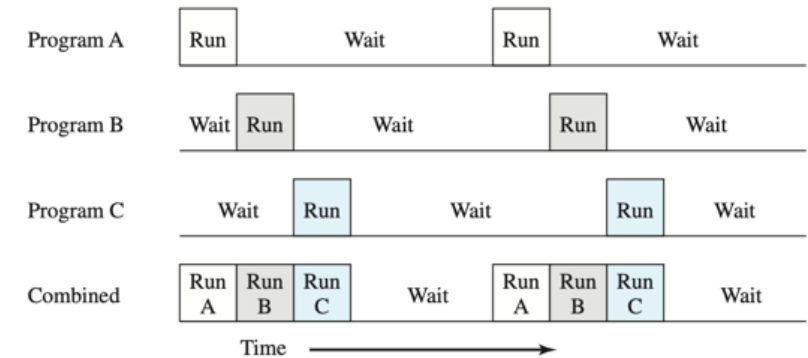
A Resource Manager

- Modern operating systems are designed to handle multiple programs (or processes) running simultaneously, a concept known as **multiprogramming** or **multitasking**.
- Goal of multiprogramming:
 - Efficient use of **CPU**, **memory**, and **I/O resources** by **interleaving** the execution of multiple programs.
- The operating system needs to share system resources between multiple processes.
 - This resource sharing can be categorized into two types: **Time sharing and Space Sharing**.

Defining Operating Systems

A Resource Manager

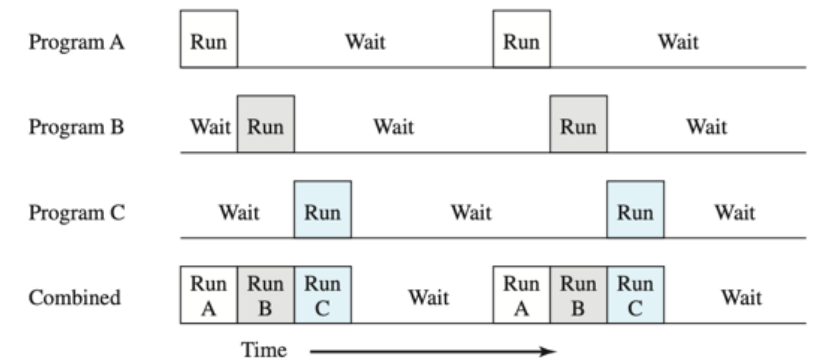
- **Time Sharing:** refers to how resources like the **CPU** or **printers** are shared by multiple processes **over time**.
- **E.g. (CPU Time Sharing or CPU scheduling):**
 - The CPU can only run one process at a time.
 - The OS uses **time slices** (small intervals of CPU time) to run each process in turn.
 - This is managed by the **OS's scheduler**, which quickly switches between processes so that each process gets CPU time and appears to run simultaneously.
- **E.g. Printers:** if multiple processes want to print, the OS manages a **print queue**, allowing each process to print in turn.



Defining Operating Systems

A Resource Manager

- Interleaving Execution of Processes
 - **Context Switching:** OS switches between processes, saving and loading their state (process control block).
- Process Scheduling:
 - OS determines the order of execution based on priority or time slices.
 - Examples: Round Robin, Priority Scheduling. (will discuss in detail next week)
- Concurrency Challenges
 - Mutual Exclusion:
 - Ensures that only one process accesses a shared resource at a time.
 - Solved by: Locks, Semaphores, Monitors.
 - Deadlock Avoidance:
 - Prevents processes from waiting indefinitely for each other. (Discuss in detail in Concurrency session)
 - Techniques: Banker's Algorithm, Deadlock Detection & Recovery. (Discuss in detail in Deadlock session)




Defining Operating Systems

A Resource Manager

- **Space sharing**

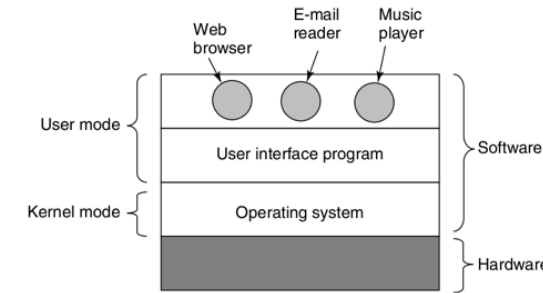
- Refers to how resources like **memory** and **disks** are shared by processes.
- The OS must **manage how much space each process uses** and ensure there is **no conflict** or overlap between them.
 - OS provides **virtual memory** to map each process's memory into physical memory locations safely.



Efficient and Restricted Execution

Kernel and User Modes in Modern Operating Systems

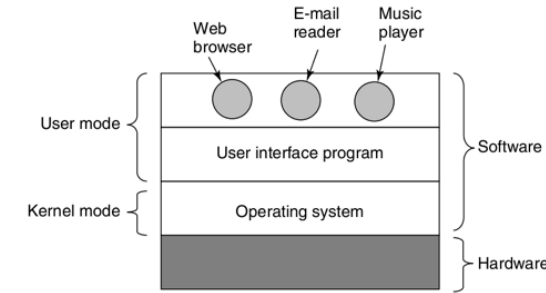
- Most computers operate in two modes to ensure **security, stability, and control** over system resources:
 - Kernel Mode:
 - **Full control over hardware:** The operating system runs with unrestricted access to all system resources.
 - **Can execute privileged instructions:** Direct access to hardware and critical system operations (e.g., managing memory, device drivers).
 - **Handles critical tasks:** Manages process scheduling, memory allocation, and I/O operations.
 - User Mode:
 - **Restricted access:** Applications (user programs) run with limited privileges.
 - **Non-privileged instructions only:** Can only perform basic operations; sensitive hardware and system resources are protected.
 - **Prevents system compromise:** Applications are restricted from performing actions that could harm the system or other processes.



Efficient and Restricted Execution

Kernel and User Modes in Modern Operating Systems

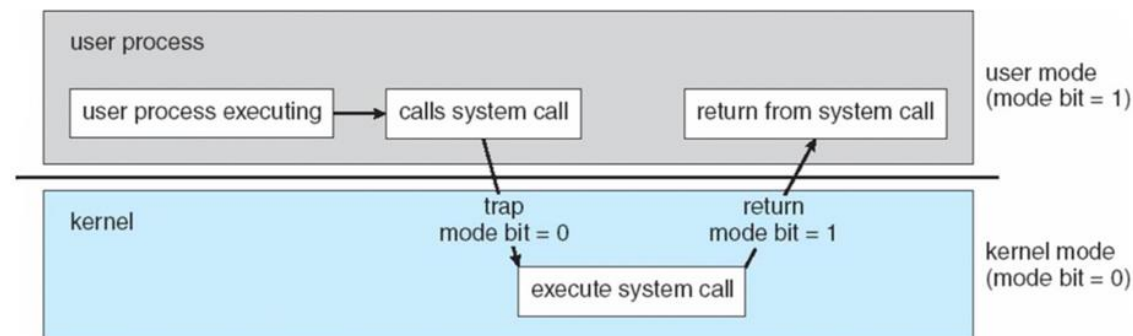
- **Privileged Instructions, e.g.**
 - I/O Operations: Direct access to devices like disks, printers.
 - Memory Management: Control over system memory allocation.
 - CPU Control: Altering process states, interrupts, or system-wide changes.
- **Non-privileged instructions (Subset of Instructions) e.g.**
 - Basic arithmetic and logical operations,
 - Instruction that accessing their own memory space,
 - Call functions that do not require hardware access.



Efficient and Restricted Execution

System call and mode switch

- **System Calls:** Applications use **system calls** or **APIs** to request privileged operations from the **Operating System (OS)**.
- **Mode Switching**
 - When an application needs to perform a privileged operation (e.g., accessing hardware), it triggers a **system call**.
 - The **OS switches** from **user mode** to **kernel mode** to execute the operation.
 - Once the operation is complete, the OS switches back to **user mode**, returning control to the application.
- **Example: Disk Read Operation**
 - When an application needs to **read from a disk**, it makes a **system call**.
 - The OS enters **kernel mode** to access the disk, performs the read operation, and then returns to **user mode**.



Process management

Call	Description
<code>pid = fork()</code>	Create a child process identical to the parent
<code>pid = waitpid(pid, &statloc, options)</code>	Wait for a child to terminate
<code>s = execve(name, argv, environp)</code>	Replace a process' core image
<code>exit(status)</code>	Terminate process execution and return status

File management

Call	Description
<code>fd = open(file, how, ...)</code>	Open a file for reading, writing, or both
<code>s = close(fd)</code>	Close an open file
<code>n = read(fd, buffer, nbytes)</code>	Read data from a file into a buffer
<code>n = write(fd, buffer, nbytes)</code>	Write data from a buffer into a file
<code>position = lseek(fd, offset, whence)</code>	Move the file pointer
<code>s = stat(name, &buf)</code>	Get a file's status information

Directory- and file-system management

Call	Description
<code>s = mkdir(name, mode)</code>	Create a new directory
<code>s = rmdir(name)</code>	Remove an empty directory
<code>s = link(name1, name2)</code>	Create a new entry, name2, pointing to name1
<code>s = unlink(name)</code>	Remove a directory entry
<code>s = mount(special, name, flag)</code>	Mount a file system
<code>s = umount(special)</code>	Unmount a file system

Miscellaneous

Call	Description
<code>s = chdir(dirname)</code>	Change the working directory
<code>s = chmod(name, mode)</code>	Change a file's protection bits
<code>s = kill(pid, signal)</code>	Send a signal to a process
<code>seconds = time(&seconds)</code>	Get the elapsed time since Jan. 1, 1970

UNIX	Win32	Description
<code>fork</code>	<code>CreateProcess</code>	Create a new process
<code>waitpid</code>	<code>WaitForSingleObject</code>	Can wait for a process to exit
<code>execve</code>	(none)	<code>CreateProcess = fork + execve</code>
<code>exit</code>	<code>ExitProcess</code>	Terminate execution
<code>open</code>	<code>CreateFile</code>	Create a file or open an existing file
<code>close</code>	<code>CloseHandle</code>	Close a file
<code>read</code>	<code>ReadFile</code>	Read data from a file
<code>write</code>	<code>WriteFile</code>	Write data to a file
<code>lseek</code>	<code>SetFilePointer</code>	Move the file pointer
<code>stat</code>	<code>GetFileAttributesEx</code>	Get various file attributes
<code>mkdir</code>	<code>CreateDirectory</code>	Create a new directory
<code>rmdir</code>	<code>RemoveDirectory</code>	Remove an empty directory
<code>link</code>	(none)	Win32 does not support links
<code>unlink</code>	<code>DeleteFile</code>	Destroy an existing file
<code>mount</code>	(none)	Win32 does not support mount
<code>umount</code>	(none)	Win32 does not support mount, so no umount
<code>chdir</code>	<code>SetCurrentDirectory</code>	Change the current working directory
<code>chmod</code>	(none)	Win32 does not support security (although NT does)
<code>kill</code>	(none)	Win32 does not support signals
<code>time</code>	<code>GetLocalTime</code>	Get the current time

What is Operating System

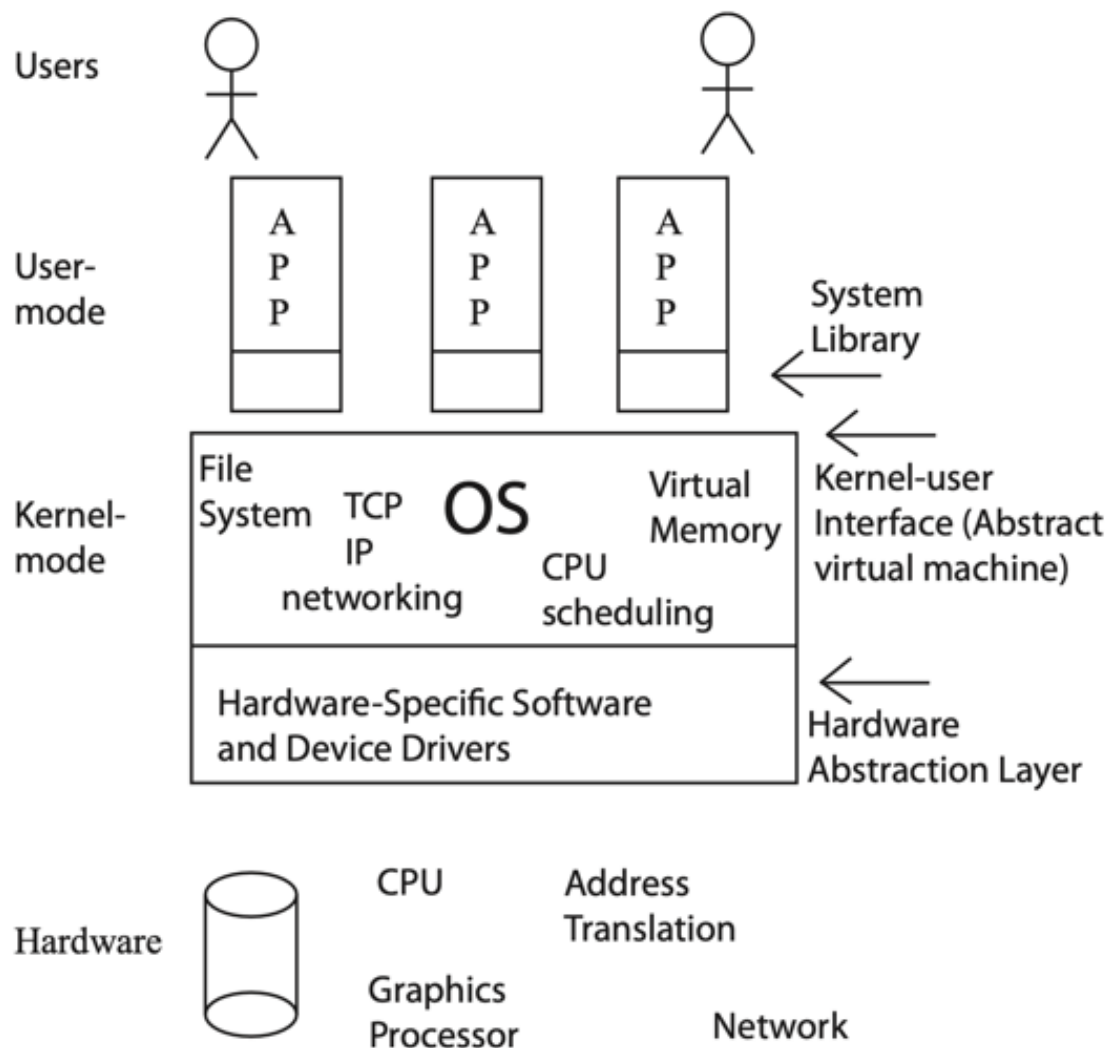
What can an OS Do for you when you write a Java program?

```
1  import java.io.FileWriter;
2  import java.io.IOException;
3  import java.io.PrintWriter;
4
5  public class Demo1 {
6      public static void main(String[] args) throws IOException {
7          FileWriter fw =
8              new FileWriter("C:/Program Files (x86)/test.txt");
9          PrintWriter pw = new PrintWriter(fw);
10         pw.close();
11     }
12 }
```

- File Creation: The OS creates the file if it doesn't exist and checks if you have permission to write to the directory.
- I/O Management: The OS manages writing data from the `PrintWriter` to the file, optimizing performance with buffering.
- Resource Allocation: The OS allocates memory for file handling and ensures that resources (like file handles) are properly released when the program finishes.
- Error Handling: The OS checks for issues like insufficient permissions or file path errors, throwing an `IOException` if something goes wrong.
- Security: The OS ensures your program only accesses files it has permission to modify.

Summary

Take-Home Message



- A virtual machine providing abstractions
 - Hide the complicate, messy, low-level hardware interface
 - Files, address, process.....
- A resource manager
 - Keep track
 - Grant resource request
 - Account for usage
 - Mediate conflicting request for different programs
- Manage in a controlled manner
 - System Calls
 - User-Kernel Mode