



**University of
Nottingham**

UK | CHINA | MALAYSIA

Operating Systems and Concurrency

Lecture 4: Process 2

**University of Nottingham, Ningbo China
2024**



Recap

Last Lecture

- Processes have “**control structures**” associated with them (process control blocks and process tables)
- Processes can have different **states** and **transition** between them (e.g. new, ready, running, blocked, terminated)
- The operating system maintains multiple **process queues** (e.g. ready queue, event queues, etc.)
- The operating system **manages processes** on the user’s behalf (e.g. fork(), exit(), . . .)

- Introduction to **process scheduling**
- Types of **process schedulers**
- **Evaluation criteria** for scheduling algorithms
- Typical **process scheduling algorithms**



Process Scheduling

Driving Factors

- The OS is responsible for **managing and scheduling** processes
- The aim of process scheduling is to assign processes to be executed by the processor or processors over time, with **good response time, throughput, and processor efficiency.**
- Process scheduling in an operating system is driven by factors such as **priority, CPU utilization, throughput, turnaround time, waiting time, response time, fairness, the nature of the process (CPU-bound vs. I/O-bound), scheduling type (preemptive vs. non-preemptive), and system load.**



- **Priority scheduling** involves **assigning different priorities to processes**, where the CPU is allocated to the process with the highest priority.
- This ensures that **critical or time-sensitive** tasks are handled promptly.
- Priorities can be assigned by **the system or user** and may be dynamic (changing based on process behavior or system state).



- The goal is to keep the **CPU as busy as possible**.
- Higher CPU utilization is generally desirable as it ensures that the processing power of the CPU is being **used efficiently**.
- Scheduling algorithms aim to reduce **idle CPU time** by optimizing process execution order.



- Throughput refers to the number of processes **completed** in a given **time frame**.
- Scheduling **should maximize throughput** by minimizing process completion time, ensuring that more processes are completed efficiently.



- Turnaround time is the **total time** taken from the **submission of a process to its completion**.
- Scheduling algorithms try to **minimize turnaround** time to improve system responsiveness and user satisfaction.
- It includes **waiting time**, **processing time**, and **any time spent waiting for I/O operations**.



- This is the total time a process spends in the **ready queue** waiting to be executed.
- **Reducing waiting time** is crucial for efficient scheduling as prolonged waits can lead to resource underutilization and poor performance.
- Scheduling strategies **aim to minimize** this for improved system performance.



- Response time measures the time from when a process is **submitted until the first response or output is produced**.
- This is important in interactive systems where users expect **immediate feedback**.
- Scheduling algorithms for interactive systems (like Round Robin) aim **to minimize response time** to ensure a responsive user experience.



- A fair scheduling algorithm ensures that **no process is starved of CPU** time and that all processes receive **equitable time for execution**.
- Fairness becomes especially important in multi-user or multi-tasking systems, where different users or applications might compete for resources.



Process Scheduling


Nature of the Process

- **CPU-Bound Processes:** processes spend most of their time performing computations on the CPU. They have **long CPU bursts**, meaning they need a lot of CPU time before requesting I/O.
 - Example: A scientific computation or a large data processing task.
- **I/O-Bound Processes:** processes spend most of their time **waiting for I/O operations to complete**, using very little CPU time. **They have short CPU bursts.**
 - Example: Reading from a file, printing documents, or interacting with a user interface.
- **Alternating Bursts:** Processes typically alternate between bursts of CPU activity (computation) and I/O requests (like reading or writing data).

Process Scheduling

Nature of the Process

- Scheduling algorithms aim to **balance** between CPU-bound and I/O-bound processes.
- **I/O-bound** processes are usually **given preference** because they spend less time on the CPU, allowing for better CPU utilization during I/O waits.





Process Scheduling

System Load

- The scheduling algorithm must adapt based on the current system load
 - (e.g., the number of processes waiting to be executed). During high loads, scheduling **might prioritize faster execution** or processes that can free resources quickly.



Process Scheduling

Classification by Time Horizon

- Process scheduling is categorized into
 - Long-term,
 - Medium-term, and
 - Short-term scheduling,
- Each with distinct roles to optimize system performance and resource utilization.



Process Scheduling

Classification by Time Horizon

- Long term scheduling:
 - Applies **to new processes** and controls the degree of multiprogramming
 - Decide which processes to admit to the system
 - A good **mix of CPU and I/O bound** processes is favorable to keep all resources as busy as possible.
 - Happens **infrequently**, as it's only triggered when a **new process is created** or when the system's load changes significantly.



Process Scheduling

Classification by Time Horizon

- Medium term scheduling:
 - Controls **swapping** and the degree of multi-programming in main memory. (Memory management)
 - Occurs when the system is **under memory pressure** or when it's necessary to adjust the degree of multiprogramming.
 - **Efficient memory utilization**, maintaining system responsiveness.



Process Scheduling

Classification by Time Horizon

- Short term scheduling (dispatcher) :
 - Manages the **ready queue**
 - Decide which process to run next
 - It is invoked in response to various events like:
 - Clock interrupts (e.g., when a time slice expires in preemptive systems),
 - I/O interrupts (e.g., when an I/O operation completes),
 - System calls (e.g., when a process makes a blocking call),
 - Semaphore operations (e.g., signaling when resources become available).
 - Invoked **very frequently**, hence must be fast
 - Minimizes response time, ensures CPU efficiency, fast dispatching.



Process Scheduling

Classification by Time Horizon (Cont.,)

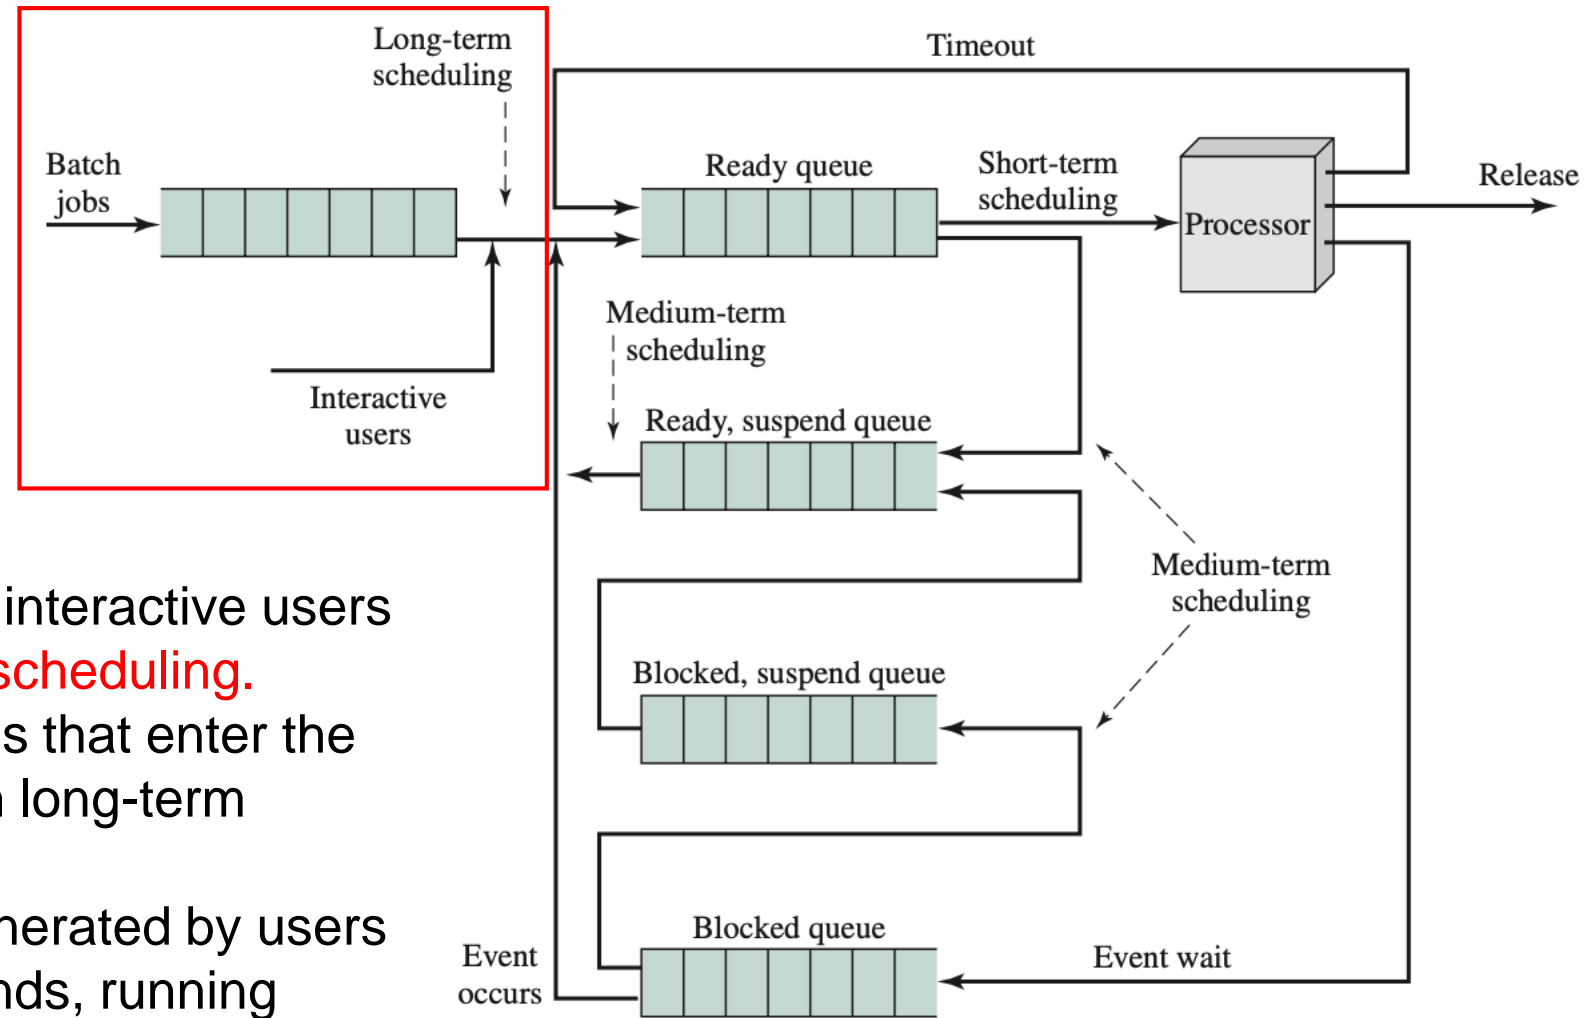


Figure. Queueing diagram for scheduling

- New processes from batch jobs and interactive users enter the system through **long-term scheduling**.
 - **Batch Jobs:** These are processes that enter the system as jobs, typically through long-term scheduling.
 - **Interactive Users:** Processes generated by users in real-time (e.g., typing commands, running interactive applications).



Process Scheduling

Classification by Time Horizon (Cont.,)

- Processes move through the **ready queue** and are selected by **the short-term scheduler** to run on the processor.

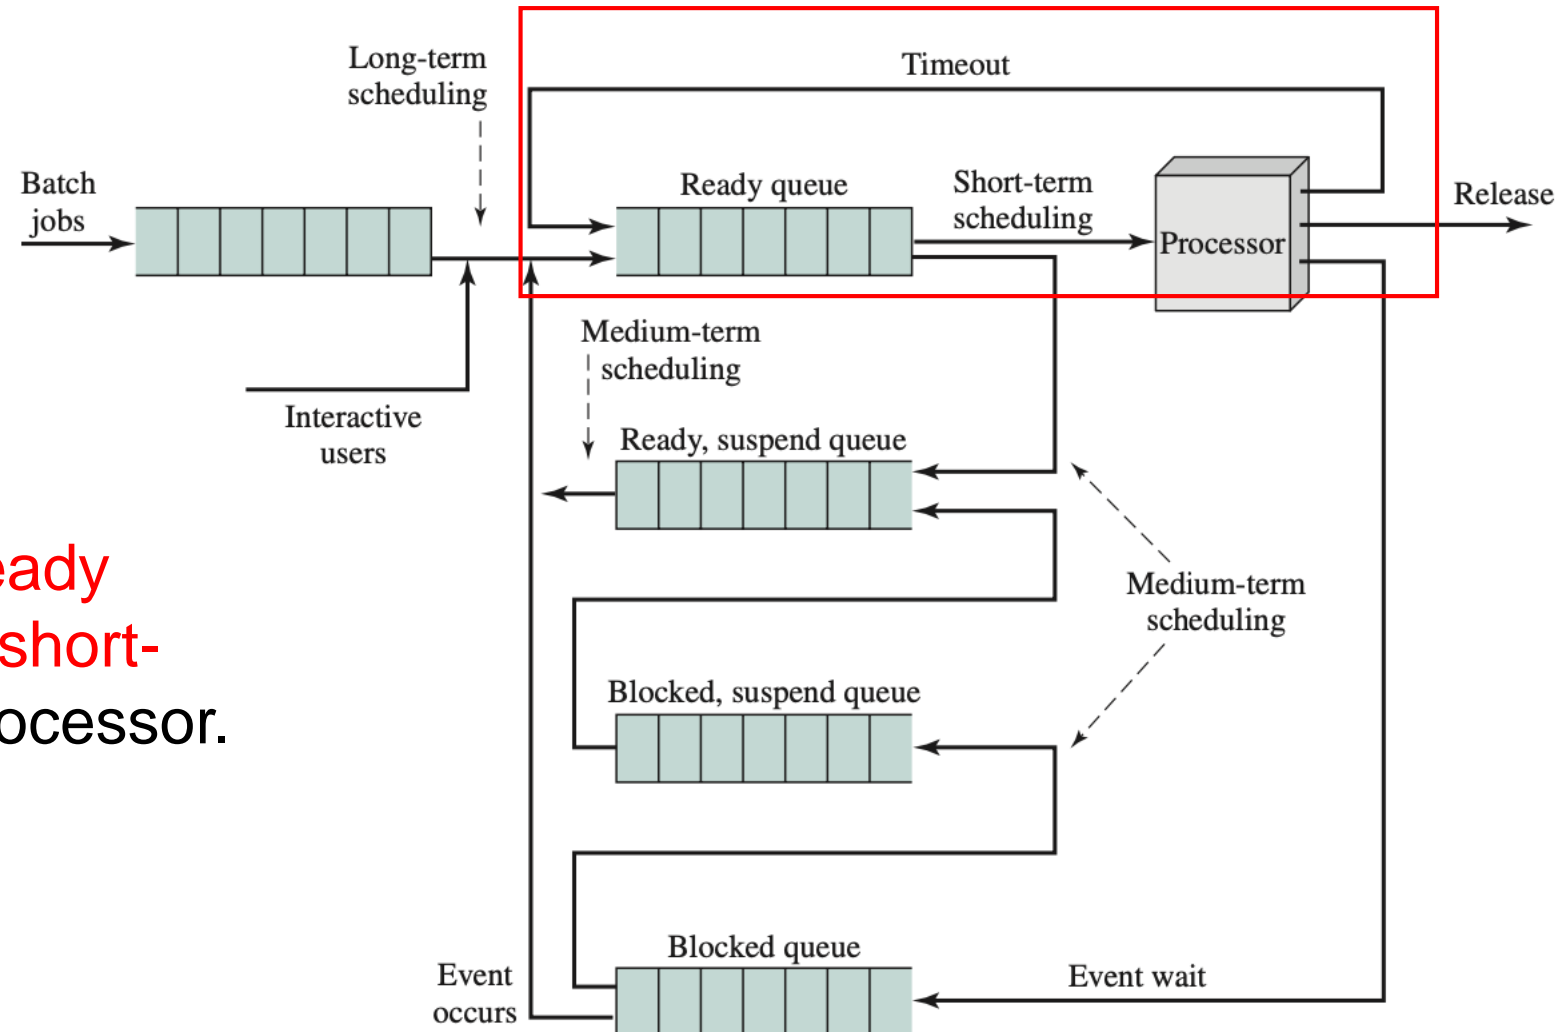


Figure. Queueing diagram for scheduling



Process Scheduling

Classification by Time Horizon (Cont.,)

- Processes may move between the **blocked, suspended, and ready queues** based on **I/O events or memory availability**, with the medium-term scheduler managing the **swapping of processes in and out of main memory**.

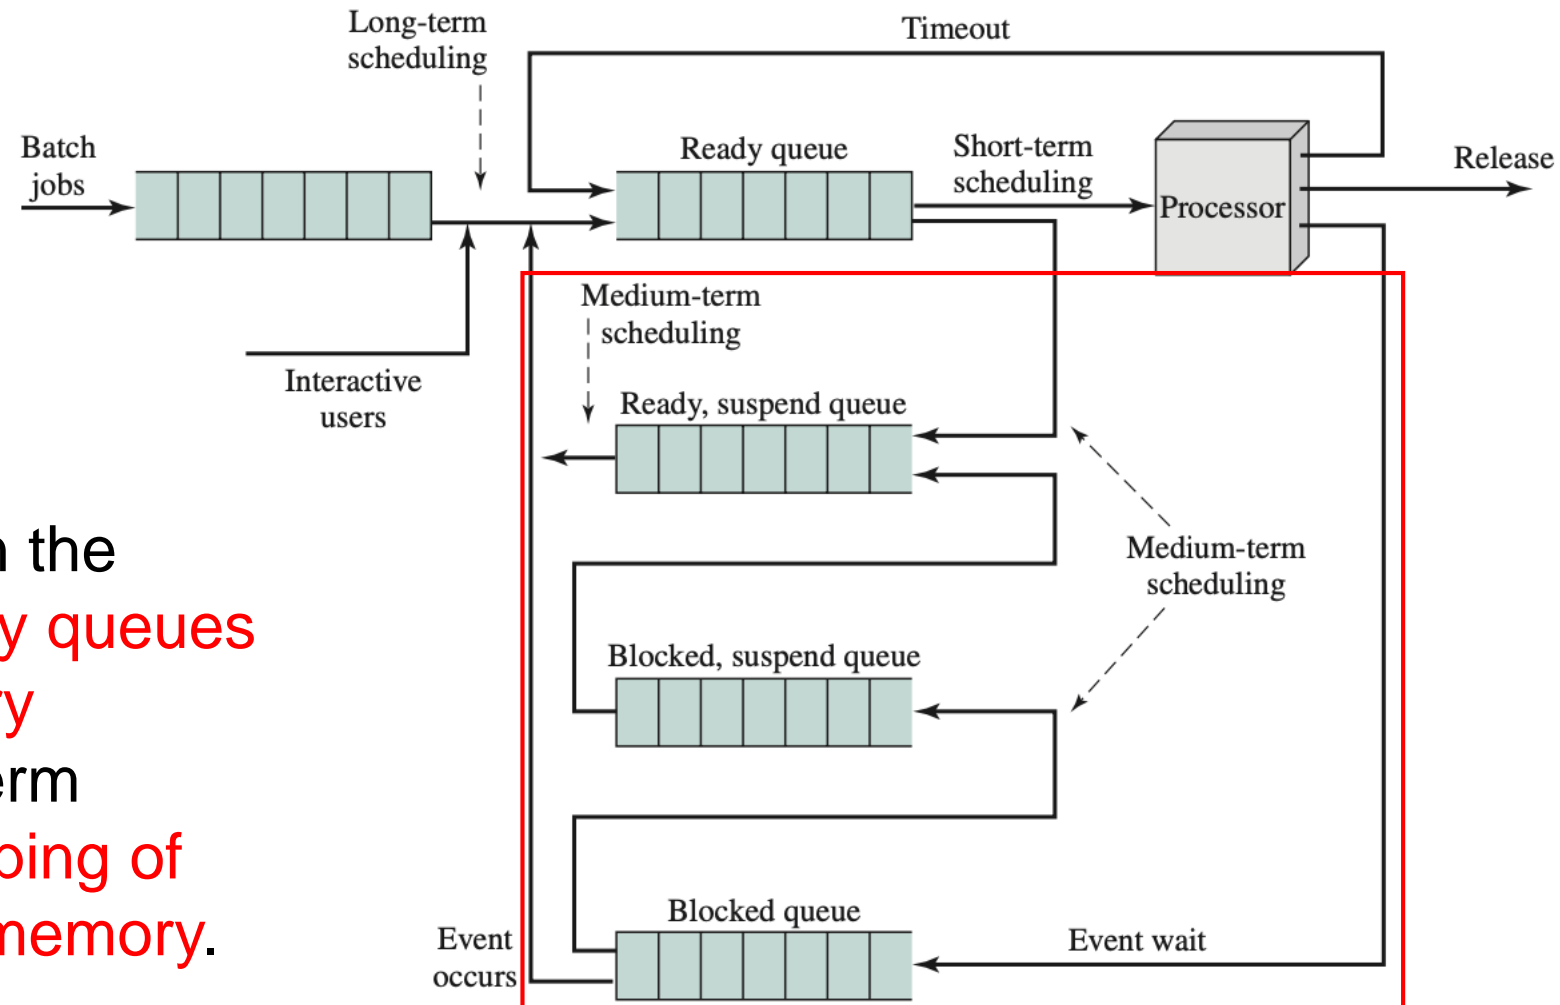


Figure. Queueing diagram for scheduling



Process Schedulers

Classification by Approach

- Process scheduling can be classified into **preemptive** and **non-preemptive** based on how processes are managed when they are running on the CPU.
 - **Non-preemptive**: processes are only interrupted **voluntarily** (e.g., I/O operation or “nice” system call – yield())
 - Windows 3.1 and DOS were non-preemptive
 - **Preemptive**: processes can be **interrupted forcefully**
 - E.g. preemptive scheduling algorithm picks a process and lets it run for a maximum of some fixed time. (clock interrupt)
 - This requires context switches, which generate **overhead**, too many of them should be avoided
 - Prevents processes from **monopolizing the CPU**
 - **Most popular** modern operating systems are preemptive.



Performance Assessment Criteria

- **User oriented criteria:**
 - **Response time:** minimize the time between creating the job and its first execution
 - **Turnaround time:** minimize the time between creating the job and finishing it
 - **Predictability:** minimize the variance in processing times
- **System oriented criteria:** (focus on **efficient utilization** of processor)
 - **Throughput:** maximize the number of jobs processed per hour
 - **Fairness:**
 - Are processing power/waiting time equally distributed?
 - Are some processes kept waiting excessively long (**starvation**)
- Evaluation criteria can be **conflicting**, i.e., **reducing the response time** may increase context switches and may **worsen the throughput** and **increase the turn around time**.



Scheduling Algorithms


Overview

- **Algorithms** considered:
 - 1 First Come First Served (**FCFS**)(**Batch System**)
 - 2 **Shortest job first (Batch System)**
 - 3 **Round Robin (Interactive System)**
 - 4 **Priority queues (Interactive System)**
- Performance measures used:
 - **Average response time**: the average of the time taken for all the processes to start
 - **Average turnaround time**: the average time taken for all the processes to finish

Scheduling Algorithms

First Come First Served

- A **non-preemptive algorithm** that operates as a **strict queueing mechanism** and schedules the processes in the same order that they were added to the queue.



- Average response time = $0 + 7 + 9 + 15 = \frac{31}{4} = 7.75$
- Average turn around time = $7 + 9 + 15 + 20 = \frac{51}{4} = 12.75$

Scheduling Algorithms

First Come First Served (Cont.,)

- Advantages: **positional fairness** and easy to implement
- Disadvantages:
 - Favours **long processes** over short ones (think of the supermarket checkout!)
 - Favours **CPU bounds processes** over I/O bounds ones.


Process	Arrival Time	Service Time (T_s)	Start Time	Finish Time	Turnaround Time (T_r)	T_r/T_s
W	0	1	0	1	1	1
X	1	100	1	101	100	1
Y	2	1	101	102	100	100
Z	3	100	102	202	199	1.99
Mean					100	26



Scheduling Algorithms

Shortest Job First

- A **non-preemptive algorithm** that starts processes in order of **ascending processing time** using a provided/known estimate of the processing
- Advantages: always result in the **optimal turn around time**
- Disadvantages:
 - **Starvation** might occur: when longer processes may never get CPU time if there are constantly shorter processes arriving.
 - **Fairness** and **predictability** are compromised:
 - SJF is not fair in a system where fairness is critical, as **short processes** are always favored over long ones, leading to **unfair treatment** of longer processes.
 - Users with long-running processes may not be able to predict when their processes will get CPU time,
 - **Requirement of Knowing Processing Times:**
 - SJF requires the processing time (or burst time) of each process to be known in advance or accurately estimated. This information is **not always readily available or easy to estimate**.




- Average response time = $0 + 2 + 7 + 13 = \frac{22}{4} = 5.5$
- Average turn around time = $2 + 7 + 13 + 20 = \frac{42}{4} = 10.5$

Scheduling Algorithms

Round Robin

- A preemptive version of FCFS that forces **context switches** at **periodic intervals** or **time slices**

- Processes run in the order that they were added to the queue
- Processes are forcefully **interrupted by the timer**
- If a time slice is only **used partially**, the next process **starts immediately**



- Average response time = $0 + 1 + 2 + 3 = \frac{6}{4} = 1.5$
- Average turn around time = $6 + 17 + 19 + 20 = \frac{62}{4} = 15.5$



Scheduling Algorithms

Round Robin

- Choosing the Right Time Slice(Quantum) : **is a critical parameter** in the Round Robin algorithm, as it affects both system responsiveness and throughput.
 - Small Time Slice:
 - **Improves response time**: A small time slice (e.g., 1 ms) ensures that no process waits long for CPU time. This is beneficial for interactive systems where quick responses are crucial.
 - **But increases overhead**: The frequent preemptions mean **more context switches**, which introduces overhead and reduces CPU efficiency.
 - Large Time Slice:
 - **Improves throughput**: A large time slice (e.g., 1000 ms) **reduces the number of context switches**, allowing processes to use the CPU for longer periods.
 - **But increases response time**: A larger time slice **increases the average time** a process has to wait before it gets CPU time, especially in systems with many processes.



Scheduling Algorithms

Round Robin (Cont.,)

- Advantages
 - **Improved Response Time:** RR ensures better responsiveness, especially in interactive systems, because processes are not starved, and each one gets CPU time in a predictable manner.
 - **Effective for General-Purpose Time-Sharing Systems:** well-suited for multi-user, time-sharing environments, as it ensures fairness.
 - Each process gets an equal opportunity to run, making it ideal for systems with multiple interactive users or tasks.
 - **Fairness:** Every process gets an equal chance at the CPU, preventing any process from monopolizing it.
 - This fairness ensures that lower-priority processes do not starve, unlike priority-based scheduling.

Scheduling Algorithms

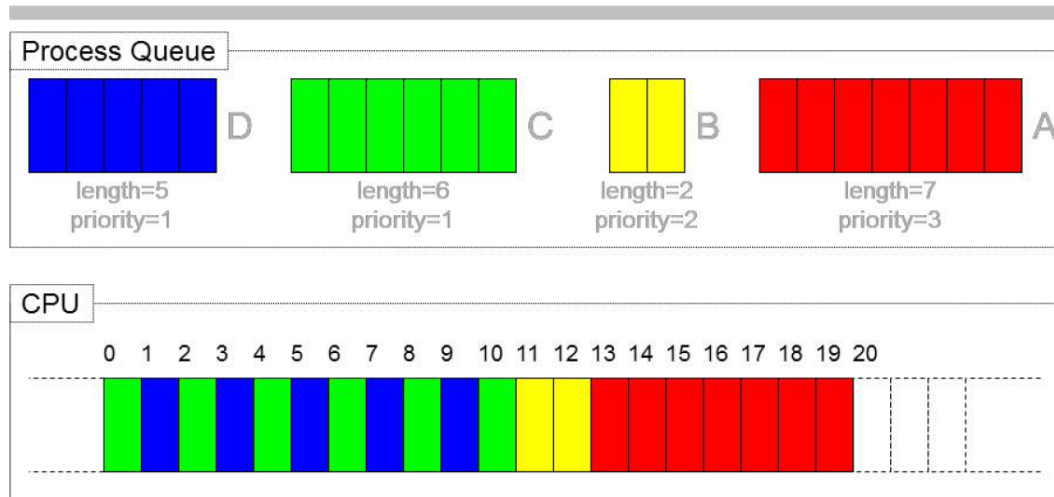
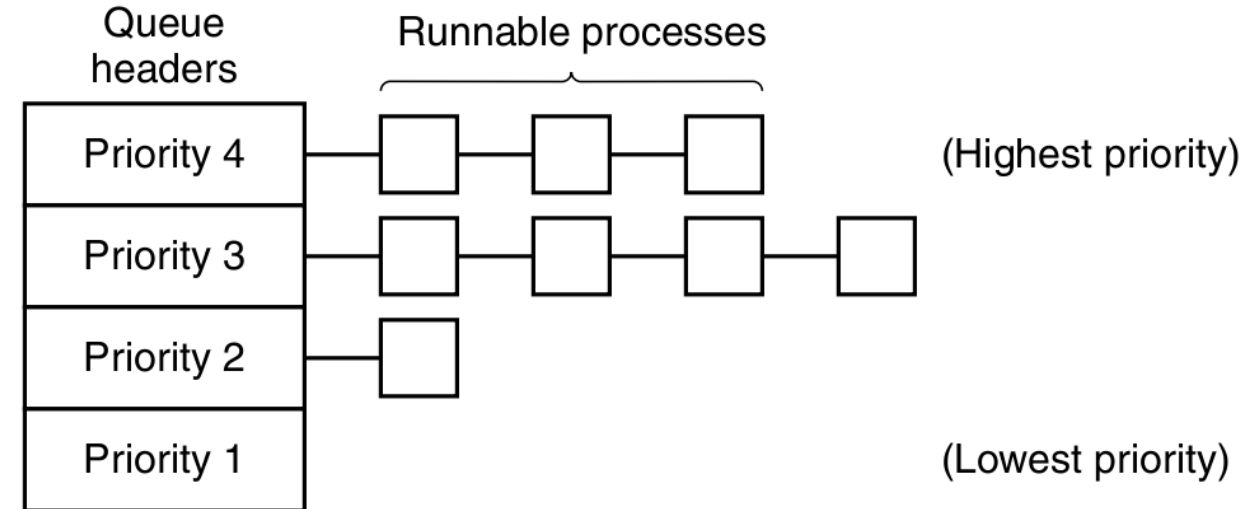
Round Robin (Cont.,)

- Disadvantage
 - **Frequent context switches** occur, especially when the time slice is short.
 - Each context switch incurs overhead (saving and restoring the state of processes), leading to reduced CPU efficiency.
 - **Favor CPU-bound processes** over I/O-bound processes.
 - Since CPU-bound processes use up their entire time slice, they are continuously moved to the end of the queue and come back for more execution. In contrast, I/O-bound processes typically require shorter bursts of CPU time and spend more time waiting for I/O. As a result, they may not get as much CPU time relative to CPU-bound processes.
 - **Can Reduce to FCFS:** If the time slice is too large, Round Robin behaves similarly to First-Come, First-Served (FCFS).
 - When the time slice is large, processes may finish within their time slice, resulting in no preemption, and the scheduler effectively **becomes non-preemptive, resembling FCFS scheduling.**

Scheduling Algorithms

Priority Queues

- A **preemptive algorithm** that schedules processes by priority (high to low)
 - The process priority is saved in the **process control block**



- Average response time = $0 + 1 + 11 + 13 = \frac{25}{4} = 6.25$
- Average turn around time = $10 + 11 + 13 + 20 = \frac{54}{4} = 13.5$



Scheduling Algorithms

Priority Queues

- Advantages:
 - Improved Responsiveness for Important Tasks:
 - ensures that important or time-sensitive processes (those with high priority) are executed as soon as possible, improving response time for critical tasks.
 - Flexibility:
 - Different processes can be treated differently based on their priority. For instance, real-time processes (e.g., handling urgent I/O) can be given higher priority over background tasks.
 - Suited for Real-Time Systems:
 - Real-time systems that need to guarantee the execution of high-priority tasks (e.g., embedded systems or multimedia applications) can benefit from priority scheduling as it allows these tasks to run without delay.



Scheduling Algorithms

Priority Queues

- Disadvantage
 - **Potential for Starvation:**
 - Low-priority processes may suffer from starvation, where they wait indefinitely if higher-priority processes continuously arrive.
 - This is particularly problematic in preemptive priority scheduling systems.
 - **Response Time Increases for Lower-Priority Processes:**
 - While high-priority tasks may benefit from low response times, lower-priority processes may have significantly longer response times.
 - This happens because every time a higher-priority process arrives, the lower-priority process is either preempted (in preemptive scheduling) or has to wait until all higher-priority processes finish.
 - **Increased Overhead in Preemptive Priority Scheduling:**
 - Frequent preemption of lower-priority processes by higher-priority ones can result in high context-switching overhead, leading to inefficiency, especially if the higher-priority processes have short CPU bursts but arrive frequently.



Summary

Take Home Message

- The OS is responsible for **process scheduling**
- Different types of schedulers exist (e.g. pre-emptive, etc.)
- Different **evaluation criteria** exist for process scheduling
- Different **algorithms** should be considered