



University of
Nottingham

UK | CHINA | MALAYSIA

Operating Systems and Concurrency

Lecture 7:
Concurrency

University of Nottingham Ningbo
China, 2024



- Threads vs. processes
- Different thread implementations
- Scheduling of thread
- POSIX Threads (PThreads)



- Concurrency Definition
 - Examples of Concurrency Problems
- Race condition
- Critical section
- Critical Section Problem Solutions



Concurrency- context

- Concurrency is the **execution of the multiple instruction sequences at the same time.**
 - It happens in the operating system when there are **several process /threads** running in **parallel or execute concurrently.**
 - The running process/threads always communicate with each other through **shared memory or message passing or can share resources** (e.g., devices, variables, memory, data structures, etc.)

- There are **several motivations** for allowing concurrent execution:
 - **Physical resource sharing** : Multiuser environment since hardware resources are limited.
 - **Logical resource sharing**: Shared file(same piece of information).
 - **Computation speedup**: Parallel execution
 - **Modularity**: Divide system functions into separation processes.



Concurrency- context

- A process/thread can be interrupted at any point in time (I/O, timer)
 - The process “state” is saved in the process control block.
 - Sharing data can lead to **inconsistencies** (e.g. when interrupted manipulating data)
 - i.e., the outcome of execution may depend on the **order in which instructions are carried out.**
- The outcome of programs may become **unpredictable.** (How?)



How? Example1-Incrementing a counter

- Let see **Counter++** statement:
 - It appears to be a single operation in **high-level programming languages** like C or Java, is actually composed of several steps at the machine level, and **these steps are not atomic**.
 - In machine language it consists of three separate actions (threads) in practice
 1. Read the value of counter and **store it in a register**
 2. Add one to the value in the register
 3. Store the value of the register **in counter**
 - i.e. they can be interrupted by, e.g., **context switches**.

Atomic operation: A function or action implemented as a sequence of one or more instructions that appears to be **indivisible**; that is, no other process can see an intermediate state or interrupt the operation. [Starlings page 201]



How? Example1-Incrementing a counter

- Consider **two processes/threads** and the following **interleaved sequence of instructions** (**they do Not interact**):
- And, assume that the initial value of counter ==1

Process 1:

...

Read counter -> register

Add 1 to value in register

Store value in counter

...

...

...

Process 2:

...

...

...

...

Read counter -> register

Add 1 to value in register

Store value in counter

- After the P1 and P2 is executed the final result of counter is **3**.



How? Example1-Incrementing a counter

- Consider **two processes/threads** and the following **interleaved sequence of instructions** (**they DO interact**):
- Assume that the initial value of counter == 1

Process 1:

...

Read counter -> register

...

Add 1 to value in register

Store value in counter

...

...

Process 2:

...

...

Read counter -> register

...

...

Add 1 to value in register

Store value in counter

- After the P1 and P2 is executed the final result of counter is **2**.



How? Example 2-Shared procedures

- Consider the following **code shared** between threads/processes
- **chin** and **chout** shared global variables

```
void print()  
{  
    chin = getchar();  
    chout = chin;  
    putchar(chout);  
}
```



How? Example 2-Shared procedures

- Consider **two processes/threads** and the following **interleaved sequence of instructions** (they do **NOT** interact):
- Assume that the input char is **X** and **Y**

Process 1:

```
...  
chin = getchar();  
chout = chin;  
putchar(chout);  
...  
...  
...
```

Process 2:

```
...  
...  
...  
...  
chin = getchar();  
chout = chin;  
putchar(chout);
```

- The output is **XY**



How? Example 2-Shared procedures

- Consider **two processes/threads** and the following **interleaved sequence of instructions** (they do interact):
- Assume that the input char is **X** and **Y**

Process 1:

```
...  
chin = getchar();  
...  
chout = chin;  
putchar(chout);  
...  
...
```

Process 2:

```
...  
...  
chin = getchar();  
...  
...  
chout = chin;  
putchar(chout);
```

- The output is **YY**



How? Example 3- Bounded Buffers

- Consider a **bounded buffer** in which N **items** can be stored
- A **counter** is maintained to count the number of items currently in the buffer
 - Incremented** when an item is **added**
 - Decrement**ed when an item is **removed**
- Similar **concurrency problems** as with the calculation of sums happen in the consumer/producer problem

```
while (true) {  
    //while buffer is full  
    while (counter == BUFFER_SIZE) ; /* do nothing */  
  
    // add item when space becomes available  
    buffer[in] = new_item;  
    in = (in + 1) % BUFFER_SIZE;  
    counter++;  
}
```

Producer

```
while (true) {  
    // wait until items in buffer  
    while (counter == 0); /* do nothing */  
  
    consumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    counter--;  
    /* consume the item in next consumed */  
}
```

Consumer



Race Conditions-Data Inconsistency (example)

```
//Counter ++
//where register1 is the one of the local
//CPU registers
register1 = counter
register1 = register1 + 1
counter = register1
```

```
// Counter --
//where register2 is the one of the local
//CPU registers
register2 = counter
register2 = register2 - 1
counter = register2
```

- Suppose we
 - Counter=5 // is a shared variable
 - Two process, Producer and Consumer.
 - Following the execution of counter++ and counter-- statements, The value of variable counter can be 4,5,6!
 - However, the **only correct result is counter=5**,
 - One of such interleaving is the following:

T ₀ :	producer	execute	register ₁ = counter	{register ₁ = 5}
T ₁ :	producer	execute	register ₁ = register ₁ + 1	{register ₁ = 6}
T ₂ :	consumer	execute	register ₂ = counter	{register ₂ = 5}
T ₃ :	consumer	execute	register ₂ = register ₂ - 1	{register ₂ = 4}
T ₄ :	producer	execute	counter = register ₁	{counter = 6}
T ₅ :	consumer	execute	counter = register ₂	{counter = 4}

- If we reversed the order of the statements at T₄ and T₅, counter==6
- A situation like this, is known as a **race condition**.



Race Conditions-Data Inconsistency

- A **race condition occurs** when multiple threads/processes **access shared data** and the result is dependent on **the order in which the instructions are interleaved**.
 - I.e., the final result depends on how the instructions are interleaved
- We will discuss **mechanisms** to provide **controlled/synchronized** access to data **and avoid race conditions**

- **A critical Section:** is a section of code within a process that requires access to **shared resources** and **that must not be executed** while another process is in a corresponding section of code.
 - Suppose n process $\{P_0, P_1, \dots, P_{n-1}\}$.
 - Each process has segment of code, **called a critical section**, in which the process may changing **common variables, updating a table, writing file, and so on**.
 - When one process is executing in its CS, no other process are executing in its critical section.

```
do
{
    . . .
    // entry to critical section
    critical section, e.g. counter++;
    // exit critical section
    remaining code
    . . .
} while (...);
```

Fig: A general structure of a typical process P_i



A Solution to the Critical Section(CS) Problem

- A **solution to the critical section** problem must satisfy the following three requirements:
 - **Mutual exclusion**: If process P_i is executing in its critical section, then no other processes **can be executing in their critical sections**.
 - **Progress**: The progress condition ensures that **if no process is currently in its critical section and some processes want to enter, then only the processes that are not in their remainder sections** (i.e., processes that want to enter the critical section) can participate in the decision of which will enter next. This decision cannot be postponed indefinitely.
 - **Fairness/bounded waiting**: There **exists a bound, or limit**, on the **number of times that other processes are allowed to enter their** critical sections after a process has made a request to enter its critical section and before that request is granted.
 - This ensures **fairness and prevents starvation**, meaning every process eventually **gets to enter the critical section**.



Real-World Example: A Shared Kitchen



[This Photo](#) by Unknown Author is licensed under [CC BY-NC-ND](#)

Imagine **a shared kitchen** in a college dormitory where multiple students want to cook their meals. This kitchen has limited resources like **stove tops**, a **refrigerator**, and **cooking utensils**.

- **Mutual Exclusion:**

- Only **one student** should be allowed to use a particular resource (e.g., a stovetop or the refrigerator) at a time.
- If two students try to use the same resource simultaneously, there might be accidents, confusion, or resource conflicts.



Real-World Example: A Shared Kitchen



[This Photo](#) by Unknown Author is licensed under [CC BY-NC-ND](#)

Imagine **a shared kitchen** in a college dormitory where multiple students want to cook their meals. This kitchen has limited resources like **stovetops**, a **refrigerator**, and **cooking utensils**.

- **Progress:**

- If the **kitchen is empty (no one is cooking)** but there are students waiting to cook, they should be able to enter the kitchen and start using the resources.
- The decision of which student gets access should be **based on those who are ready and waiting**, not on arbitrary factors or students who are not yet ready to cook.
- A fair system could be implemented, such as a waiting list or a **"first-come, first-served"** queue, ensuring that the students who are actively ready to cook are the ones deciding who goes next, without unnecessary delays or indecision.



Real-World Example: A Shared Kitchen



[This Photo](#) by Unknown Author is licensed under [CC BY-NC-ND](#)

Imagine **a shared kitchen** in a college dormitory where multiple students want to cook their meals. This kitchen has limited resources like **stovetops**, a **refrigerator**, and **cooking utensils**.

- **Fairness (Bounded Waiting):** Suppose a student is waiting to use the stovetop. **There should be a limit to how many other students can finish cooking before this student gets a chance to cook.** This ensures that no student has to wait indefinitely while others keep cooking.

Example of CS

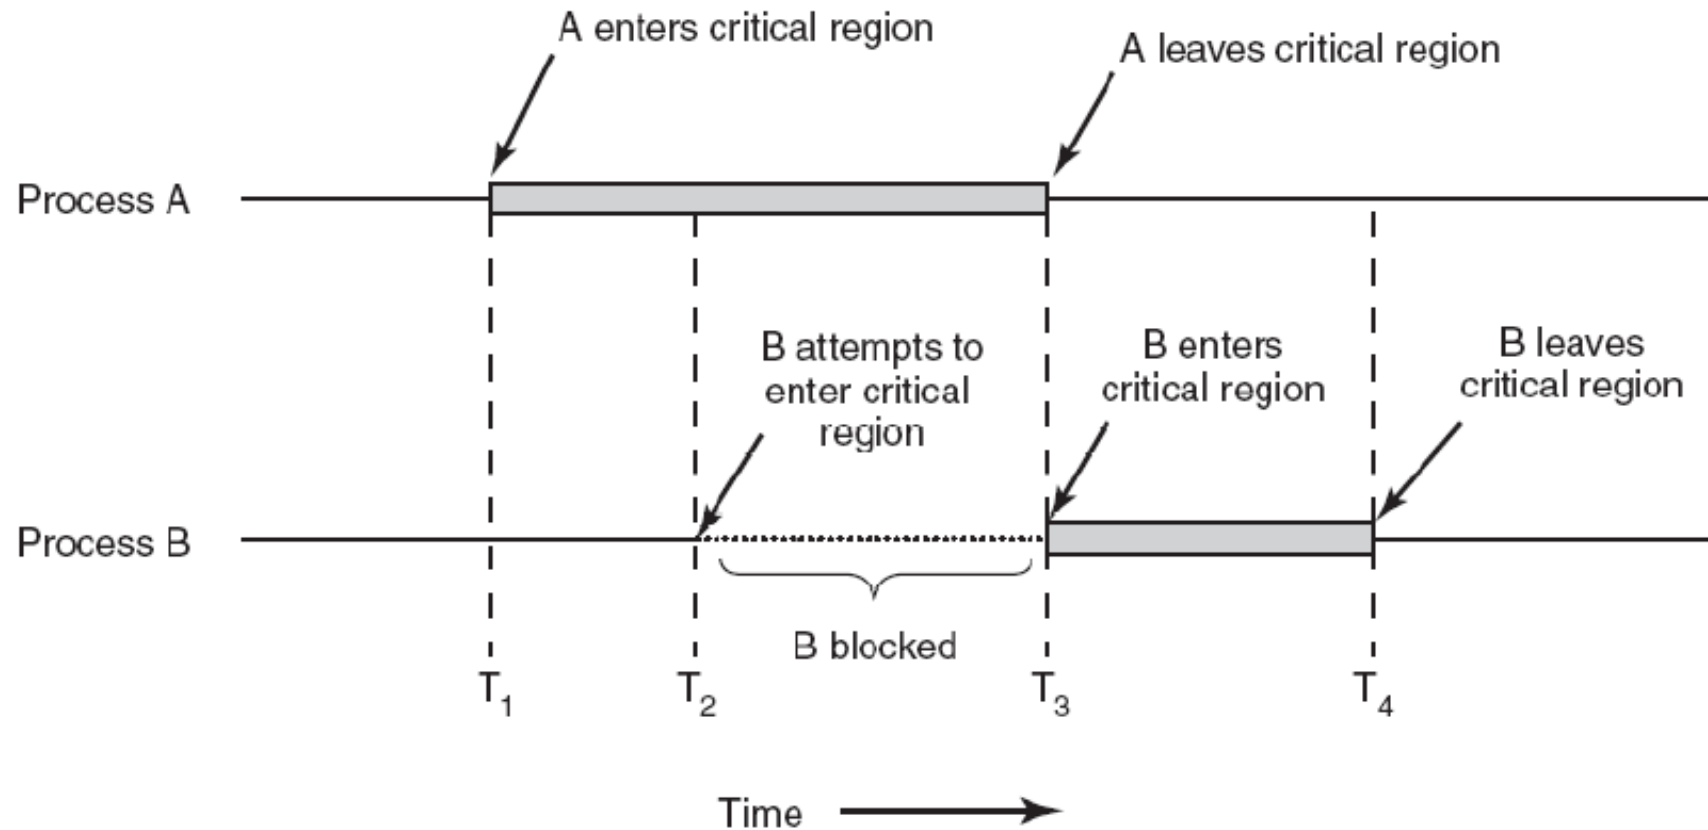


Figure: Tanenbaum, 4th ed., section 2.3.2



Approaches Handling the critical section problem

- The approaches to handle the critical section problem, categorized by:
 - Software based
 - Hardware based
 - Operating System approach



Approaches Handling the critical section problem

Software based

- These involve implementing synchronization mechanisms directly within the application code **using variables, flags, and custom algorithms**.
- These solutions work at the **user-space** level, where **no special hardware support** or **operating system intervention** is needed.
 - Examples:
 - **Peterson's Solution**,
 - Lamport's Bakery Algorithm,
 - Dekker's Algorithm and Szymanski's Algorithm:

Approaches Handling the critical section problem

Hardware based

- These approaches rely on **low-level hardware instructions** to provide mutual exclusion.
- They ensure that certain operations (like reading and writing shared memory) are atomic, meaning that they **cannot be interrupted and will always complete fully**.
- Examples:
 - Disabling interrupt,
 - Test_and_set() and
 - Compare_and_swap()



Approaches Handling the critical section problem

Operating System approach

- Provides synchronization primitives to manage access to critical sections.
- These primitives are implemented in the kernel and provide efficient control over shared resources, often relying on hardware mechanisms internally.
 - Examples:
 - Mutexes
 - Semaphores
 - Monitors



Concurrency within the OS-Data Structures

- Two general approaches that operating systems use to **handle critical sections**, particularly in kernel mode: Preemptive kernels and Non-preemptive kernels
- Preemptive kernels:
 - A process in kernel mode **can be interrupted and another process can be scheduled**.
 - **Multiple processes** can be active in the kernel at once, leading to potential **race conditions**.
 - Requires synchronization:
 - Uses locks, mutexes, and atomic operations to protect shared data.
 - E.g. Windows and Linux



Concurrency within the OS-Data Structures

- Two general approaches that operating systems use to **handle critical sections**, particularly in kernel mode: Preemptive kernels and Non-preemptive kernels
- Non-preemptive kernels
 - A process runs in kernel mode until it finishes, blocks, or yields.
 - **No race conditions:** Only one process is active in the kernel at a time, simplifying design.
 - E.g. Older or simpler OS, such as early versions of **MS-DOS or embedded systems**.



University of
Nottingham

UK | CHINA | MALAYSIA

Quiz



- **Problems with synchronization** and concurrent/parallel code
- Concurrency from an **OS perspective**
- Concept of **mutual exclusion**
- Requirements and **approaches** for mutual exclusion

- Modern Operating Systems(Tanenbaum):**Chapter2(2.3)**
- Operating System Concepts(Silberschatz):**Chapter5(5.1-2)**
- Operating Systems: Internals and Design Principles
(Starlings):**Chapter5(5.1-2)**