



University of
Nottingham

UK | CHINA | MALAYSIA

Operating Systems and Concurrency

Lecture 12:
Concurrency

University of Nottingham,
Ningbo China, 2024



Recap Last Lecture

- The Multiple Consumer, Multiple Producer, Bounded Buffer
- The dining philosophers problem
 - Scenarios
 - Solutions



Today's class

- What are the monitors and how can we use them?
- Monitor Semantics/Structure
- limitation
- Monitor Implementation
- Solving synchronization problems with monitors
 - Pseudocode Example: Monitor for producer consumer Bounded Buffer
 - The dining philosophers problem Solution with Monitor
- Semaphores versus Condition Variables



What's wrong with Semaphores?

- Global Variables: Semaphores are shared variables accessible from anywhere in the program, **increasing the risk of misuse**.
- No Connection to Data: Lack of inherent linkage between the semaphore and the resource it protects, leading to confusion in complex systems.
- Unrestricted Access: Any part of the code can signal or wait on semaphores, making it hard to enforce correct usage and **increasing the chance of errors**.
- Dual Purpose Functionality: Serve both mutual exclusion and scheduling, **complicating their management and leading to potential misuse**.
- No Enforcement of Proper Usage: No built-in checks ensure that semaphores are used correctly, **leading to risks like deadlock**.



Monitors

- **Deadlock:** Occur when two tasks try to lock two different semaphores in a different order.
 - E.g. if the programmer mistakenly develop this


```
signal(mutex);  
...  
critical section  
...  
wait(mutex);
```

Suppose that a process interchanges the order in which the wait() and signal() operations.

```
wait(mutex);  
...  
critical section  
...  
wait(mutex);
```

Suppose that a process replaces signal(mutex) with wait(mutex)

- Solution: use a higher level primitive called **monitors**.




Sir Charles Antony Richard Hoare

- Hoare 1974
- Monitor is an **abstract data type** designed for handling and defining shared resources in concurrent programming.
- A monitor is similar to a **class** that ties the **data**, **operations**, and in particular, the **synchronization operations** all together,
- Unlike classes,
 - monitors guarantee **mutual exclusion**, i.e., only one thread may execute a given monitor method at a time.
 - monitors **require all data to be private**.


Monitors: Components

- Shared Private Data
 - Represents the resource being managed.
 - This data **cannot be accessed directly from outside** the monitor, ensuring encapsulation and protection.
- Procedures that operate on the data
 - These are the **gateway to accessing and manipulating the shared data**.
 - Procedures can only operate on the data **local to the monitor, promoting safe access patterns**.
 - The monitors can have **N procedures**
- Synchronization primitives:
 - Monitors include mechanisms to synchronize access among threads that interact with the procedures.
 - This prevents **race conditions and ensures mutual exclusion** when accessing shared resources.



Monitor Monitor Semantics

- Monitors guarantee **mutual exclusion using locks**.
 - Only one thread can execute a monitor procedure at any time.
 - “in the monitor”
- If **second thread invokes** a monitor procedure at that time
 - It will block and wait for entry to the monitor
 - Need for a wait queue





Monitor Condition Variables

- How can we change `pop()` to wait until something is on the queue?
 - Logically, we want to go to sleep inside of the critical section.
 - But if we hold on to the lock and sleep, then other threads cannot access the shared queue, add an item to it, and wake up the sleeping thread
 - => The thread could sleep forever
- **Solution:** use condition variables
 - Condition variables enable a thread to sleep inside a critical section/Monitor.
 - Any lock held by the thread is atomically released when the thread is put to sleep

```
Monitor monitor_name
{
    // shared variable declarations

    procedure P1(. . . .) {
        . . . . .
    }

    procedure P2(. . . .) {
        . . . . .
    }
    :
    procedure PN(. . . .) {
        . . . . .
    }

    initialization_code(. . . .) {
        . . . . .
    }
}
```

For example:

```
Monitor stack
{
    int top;
    void push(any_t *) {
        . . . . .
    }


    any_t * pop() {
        . . . . .
    }

    initialization_code() {
        . . . . .
    }
}
```

Monitors

Synchronization Using Monitors


- Defines Condition Variables:
 - Condition variables are used to allow threads to wait for certain conditions to be met.
 - select a meaningful name (good programming practice)
 - condition x;
- 3 atomic operations on Condition Variables
 - **x.wait(cond_var_name)**: release monitor lock, sleep until woken up
 - Condition variables have a waiting queue
 - **x.notify/signal(cond_var_name)**: wake one process waiting on condition.
 - Notifier gives up lock and waiter runs immediately (if there is one)
 - **x.notifyAll/signalall(cond_var_name)**: wake all processes waiting on condition
 - Useful for resource manager
- **Rule**: thread must **hold the lock** when doing condition variable operations.



Monitors

Types of wait queues

- Monitors have two kinds of “wait” queues
 - Entry to the monitor:** has a queue of threads waiting to obtain mutual exclusion so they can enter
 - Condition variables:** each condition variable has a queue of threads waiting on the associated condition





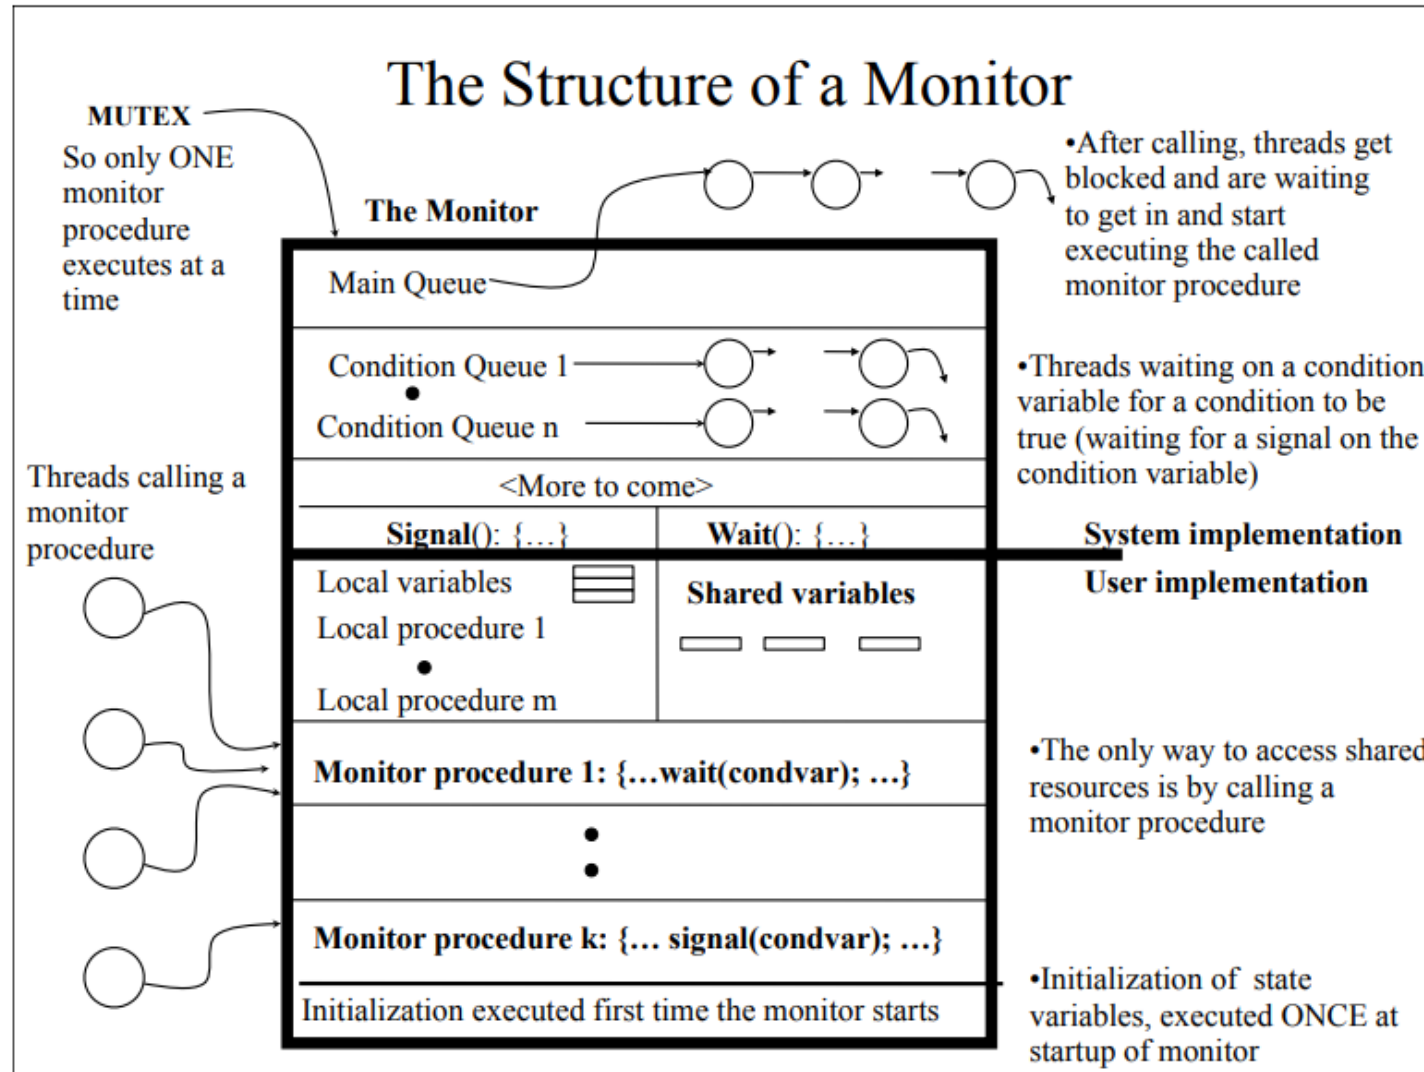
Monitor

Mesa versus Hoare Monitors

- What should happen when notify() is called?
 - If there is no waiting threads => the notifier continues and the signal is effectively lost (unlike what happens with semaphores).
 - If there is a waiting thread, one of the threads starts executing, others must wait
- **Mesa-style: (Nachos, Java, and most real operating systems(Linux, Windows, and macOS))**
 - The thread that notifies keeps the lock (and thus the processor).
 - The waiting thread waits for the lock.
- **Hoare-style: (most textbooks)**
 - The thread that notifies gives up the lock and the waiting thread gets the lock.
 - When the thread that was waiting and is now **executing exits or waits again**, it releases the lock back to the signaling thread.

Monitor

Summary: The Structure of a Monitor





Monitor limitations

- Nested Monitors: Using one monitor within another can lead to deadlocks.
- Priority Inversion: Monitors do not address **priority inversion**, where a low-priority thread holding a monitor prevents a higher-priority thread from proceeding.
- Limited Flexibility: Monitors are **less flexible than semaphores for certain complex synchronization needs**.




Monitors: Implementation

- Monitors are a **higher-level synchronization** construct that requires deep integration with the language's concurrency model.
 - Languages with Monitor Support: **Java, Python (with limitations), C#, and Ada.**
 - Languages without Monitor Support: C, C++, JavaScript, Go, and others.

Monitors

Pseudocode Example: Monitor for producer consumer Bounded Buffer

Bounded Buffer Monitor



Rules for the buffer B:

- No Get when empty
- No Put when full
- B shared, so must have mutex between Put and Get

One condition variable for each condition:

- nonempty
- nonfull
- MUTEX is already provided by the monitor

```
/*Local functions, variables*/
int in, out;
/*Shared variable*/
int B(0..n-1), count;
/*Condition variable*/
Condition nonfull, nonempty;
```

Put (int m):

```
{ if (count==n) wait (nonfull);
  B(in):=m;
  in:=in+1 MOD n; /* MOD is % */
  count++;
  signal (nonempty); }
```

int Get:

```
{ if (count==0) wait (nonempty);
  Get:=B(out);
  out:=out+1 MOD n;
  count--;
  signal (nonfull); }
```

/* Initialization code*/

```
in:=out:=count:=0;
nonfull, nonempty:=EMPTY;
```



Monitors

Dining-Philosophers Solution Using Monitors

- Let's illustrate monitor concepts by presenting a **deadlock-free** solution to the dining philosophers problem.
- This solution imposes the restriction that a philosopher may pick up his chopsticks **only if both of the are available**.
- To code this solution,
 - we need to distinguish among three states in which we may find a philosopher. For this purpose, we introduce the following data structure:
 - `enum{thinking,hungry,eating} state[5];`
 - Philosopher i can set the variable `state[i]=eating` only if his two neighbors are not eating:
 - `(state [(i+4) % 5] !=eating (right) and (state [(i+1)%5]!=eating (left)));`
 - We also need to declare `condition self[5]` where philosopher i can **delay** himself when **he is hungry but is unable to obtain the chopsticks he needs**.

Monitors

Dining-Philosophers Solution Using Monitors

```

1  monitor DiningPhilosophers
2  {
3      enum {THINKING, HUNGRY, EATING} state[5];
4      condition self[5];
5      void pickup(int i) {
6          state[i] = HUNGRY;
7          test(i);
8          if (state[i] != EATING)
9              self[i].wait();
10     }
11     void putdown(int i) {
12         state[i] = THINKING;
13         test((i + 4) % 5);
14         test((i + 1) % 5);
15     }
16     void test(int i) {
17         if ((state[(i + 4) % 5] != EATING) &&
18             (state[i] == HUNGRY) &&
19             (state[(i + 1) % 5] != EATING)) {
20             state[i] = EATING;
21             self[i].signal();
22         }
23     }
24     initialization code() {
25         for (int i = 0; i < 5; i++)
26             state[i] = THINKING;
27     }
28 }

```

P0	P1	P2	P3	P4
T	T	T	T	T

State[N]

```

1  int main()
2  {
3
4      DiningPhilosophers.pickup(i);
5      ...
6      eat
7      ...
8      DiningPhilosophers.putdown(i);
9
10     return 0;
11 }

```

Fig: A monitor solution to the dining-philosopher problem



Monitors

Dining-Philosophers Solution Using Monitors

```
1  monitor DiningPhilosophers
2  {
3      enum {THINKING, HUNGRY, EATING} state[5];
4      condition self[5];
5      void pickup(int i) {
6          → state[i] = HUNGRY;
7              test(i);
8              if (state[i] != EATING)
9                  self[i].wait();
10     }
11     void putdown(int i) {
12         state[i] = THINKING;
13         test((i + 4) % 5);
14         test((i + 1) % 5);
15     }
16     void test(int i) {
17         if ((state[(i + 4) % 5] != EATING) &&
18             (state[i] == HUNGRY) &&
19             (state[(i + 1) % 5] != EATING)) {
20             state[i] = EATING;
21             self[i].signal();
22         }
23     }
24     initialization code() {
25         for (int i = 0; i < 5; i++)
26             state[i] = THINKING;
27     }
28 }
```

P0	P1	P2	P3	P4
T	T	H	T	T

State[N]

```
1  int main()
2  {
3
4      → DiningPhilosophers.pickup(i);
5      ...
6      eat
7      ...
8      DiningPhilosophers.putdown(i);
9
10     return 0;
11 }
```

Fig: A monitor solution to the dining-philosopher problem

Monitors

Dining-Philosophers Solution Using Monitors

```

1  monitor DiningPhilosophers
2  {
3      enum {THINKING, HUNGRY, EATING} state[5];
4      condition self[5];
5      void pickup(int i) {
6          state[i] = HUNGRY;
7          → test(i);
8          if (state[i] != EATING)
9              self[i].wait();
10     }
11     void putdown(int i) {
12         state[i] = THINKING;
13         test((i + 4) % 5);
14         test((i + 1) % 5);
15     }
16     void test(int i) {
17         if ((state[(i + 4) % 5] != EATING) &&
18             (state[i] == HUNGRY) &&
19             (state[(i + 1) % 5] != EATING)) {
20             state[i] = EATING;
21             self[i].signal();
22         }
23     }
24     initialization code() {
25         for (int i = 0; i < 5; i++)
26             state[i] = THINKING;
27     }
28 }

```

P0	P1	P2	P3	P4
T	T	E	T	T

State[N]

```

1  int main()
2  {
3
4  → DiningPhilosophers.pickup(i);
5      ...
6      eat
7      ...
8      DiningPhilosophers.putdown(i);
9
10     return 0;
11 }

```

Fig: A monitor solution to the dining-philosopher problem

Monitors

Dining-Philosophers Solution Using Monitors

```

1  monitor DiningPhilosophers
2  {
3      enum {THINKING, HUNGRY, EATING} state[5];
4      condition self[5];
5      void pickup(int i) {
6          state[i] = HUNGRY;
7          test(i);
8          if (state[i] != EATING)
9              → self[i].wait();
10     }
11     void putdown(int i) {
12         state[i] = THINKING;
13         test((i + 4) % 5);
14         test((i + 1) % 5);
15     }
16     void test(int i) {
17         if ((state[(i + 4) % 5] != EATING) &&
18             (state[i] == HUNGRY) &&
19             (state[(i + 1) % 5] != EATING)) {
20             state[i] = EATING;
21             self[i].signal();
22         }
23     }
24     initialization code() {
25         for (int i = 0; i < 5; i++)
26             state[i] = THINKING;
27     }
28 }

```

P0	P1	P2	P3	P4
T	T	E	H	T

State[N]

```

1  int main()
2  {
3
4  → DiningPhilosophers.pickup(i);
5      ...
6      eat
7      ...
8      DiningPhilosophers.putdown(i);
9
10     return 0;
11 }

```

If P3 try to eat
it will be
suspended

Fig: A monitor solution to the dining-philosopher problem

Monitors

Dining-Philosophers Solution Using Monitors

```

1  monitor DiningPhilosophers
2  {
3      enum {THINKING, HUNGRY, EATING} state[5];
4      condition self[5];
5      void pickup(int i) {
6          state[i] = HUNGRY;
7          test(i);
8          if (state[i] != EATING)
9              → self[i].wait();
10     }
11     void putdown(int i) {
12         state[i] = THINKING;
13         test((i + 4) % 5);
14         test((i + 1) % 5);
15     }
16     void test(int i) {
17         if ((state[(i + 4) % 5] != EATING) &&
18             (state[i] == HUNGRY) &&
19             (state[(i + 1) % 5] != EATING)) {
20             state[i] = EATING;
21             self[i].signal();
22         }
23     }
24     initialization code() {
25         for (int i = 0; i < 5; i++)
26             state[i] = THINKING;
27     }
28 }

```

P0	P1	P2	P3	P4
T	T	T	E	T

State[N]

```

1  int main()
2  {
3
4      DiningPhilosophers.pickup(i);
5      ...
6      eat
7      ...
8      DiningPhilosophers.putdown(i);
9
10     return 0;
11 }

```

P3 wakes up when
P2 finished eating.

Fig: A monitor solution to the dining-philosopher problem



Monitor

Semaphores versus Condition Variables

- Condition variables **do not have any history**, but semaphores do.
 - On a condition variable,
 - Signal, If no one is waiting, **the signal is a no-op.**
 - If a thread then does a **condition Wait**, it waits.
 - On a semaphore,
 - Signal, If no one is waiting, **the value of the semaphore is incremented.**
 - If a thread then does a semaphore Wait, then value is decremented of semaphore. If it is less than zero blocked else continues.
- Semaphore Wait and Signal are commutative, the result is the same regardless of the order of execution.
- Condition variables are not, and as a result they must be in a critical section to access shared variables and do their job.
- It is possible **to implement monitors with semaphores**



University of
Nottingham

UK | CHINA | MALAYSIA

Quiz



Recap Take-Home Message

- Modern Operating Systems (Tanenbaum): **Chapter 2(2.3.5, 2.5.1)**
- Operating System Concepts (Silberschatz ninth Edition): **Chapter 5(5.8)**
- Operating Systems: Internals and Design Principles (Starlings): **Chapter 5(5.3, 5.6)**