



University of
Nottingham

UK | CHINA | MALAYSIA

Operating Systems and Concurrency

Lecture 14: Deadlocks (cont.)

University of Nottingham,
Ningbo China, 2024



- Definition of deadlocks
- Deadlock occurrence
 - Minimum conditions
- Approaches to dealing with deadlocks
 - Deadlock detection algorithms
 - Deadlock recovery approaches



Overview

Today class

- Deadlock Prevention
- Deadlock Avoidance
 - Banker algorithm



Deadlock Prevention

Restrain one of the following four conditions

- Deadlock Prevention Restrain one of the following four conditions:
 - Mutual Exclusion
 - Hold and Wait
 - No Preemption
 - Circular Wait



Deadlock Prevention

Restrain one of the four conditions

- Deadlock Prevention Restrain one of the following four conditions:
 - Mutual Exclusion:
 - **Prevention Strategy:** If resources can be shared (e.g., read-only files), they don't require mutual exclusion.
 - However, certain resources (like printers or files that need write access) inherently require mutual exclusion, **so this approach doesn't work for all resources.**
 - **Limitation:** This condition cannot be fully eliminated for **non-sharable resources**, so it's **only partially effective in deadlock prevention.**



Deadlock Prevention

Restrain one of the four conditions

- Deadlock Prevention Restrain one of the following four conditions:
 - **Hold and Wait:**
 - Description: Prevention Strategy:
 - **Option 1:** Require processes **to request and be allocated all required** resources before they begin execution.
 - Pros: Simple to implement.
 - Cons: Leads to low resource utilization since resources **may be held for longer than necessary.**
 - **Option 2:** Allow processes to request resources only when they **currently hold none.**
 - Pros: Reduces the risk of deadlock.
 - Cons: **Starvation may occur**, as processes might wait indefinitely if they're repeatedly unable to acquire all resources needed.
 - **Limitation:** Both methods have practical drawbacks, **like low efficiency or potential starvation.**



Deadlock Prevention

Restrain one of the four conditions

- Deadlock Prevention Restrain one of the following four conditions:
 - **No Preemption:**
 - Prevention Strategy:
 - If a process holding resources requests additional resources and those resources **are not immediately available, the process must release all currently held resources.**
 - Example: If Process P_1 requests a resource R_1 currently held by Process P_2 , which is waiting for another resource R_2 , R_1 could be reallocated to P_1 .



- **Pros:** Helps prevent indefinite waiting and potential deadlocks.
- **Cons:** May result in **process starvation or significant process rollback** (restarting from scratch), as processes may be forced to release resources frequently.




Deadlock Prevention

Restrain one of the four conditions

- Deadlock Prevention Restrain one of the following four conditions:
 - Circular Wait:**
 - Prevention Strategy: **Impose a strict ordering of resources** and require that each process requests resources in a specific increasing order.
 - Example: If there are resources R_1, R_2, \dots, R_n each process must request resources in ascending order (e.g., first R_1 , then R_2 , and so on).



- Pros:** Effectively eliminates circular waiting, breaking the cycle required for deadlock.
- Cons:** Processes may need to be restructured to follow the ordering, which **can complicate programming** and may **not always align with the logical flow of tasks**.



- **Dijkstra's Banker's Algorithm**, developed by Edsger Dijkstra in 1965.
- The name “Banker’s Algorithm” comes from its analogy to a banking system, where a **banker must ensure they retain enough funds to satisfy the needs of all clients while processing loan requests.**
- It is modeled on the way a small-town banker might deal with a group of customers to whom he has granted lines of credit. (Years ago, banks did not lend money unless they knew they could be repaid.)





Deadlock Avoidance

Key Points of the Banker's Algorithm in this Banking Analogy

1. **Loan limits (maximum needs)** help the bank gauge worst-case demands and manage its cash flow.
2. The **safe state is crucial**: Before lending, the bank checks if it can meet all customers' maximum demands safely to avoid a shortage.
3. **Sequential repayment** and reallocation of funds enable the bank to serve multiple customers over time, ensuring no one is denied due to insufficient funds.



Banker's Algorithm (BA)

Example: Banking Analogy

What the algorithm does is check to see if granting the request leads to an unsafe state. If so, the request is denied. If granting the request leads to a safe state, it is carried out.

Has Max		
A	0	6
B	0	5
C	0	4
D	0	7

Free: 10

(a)

This state is safe

Has Max		
A	1	6
B	1	5
C	2	4
D	4	7

Free: 2

(b)

This **state is safe** because with two units left, the banker can delay any requests except C's, thus letting C finish and release all four of his resources. With four units in hand, the banker can let either D or B have the necessary units, and so on.

Has Max		
A	1	6
B	2	5
C	2	4
D	4	7

Free: 1

(c)

Consider what would happen if a request from B for **one more unit** were granted in (b). We would have situation (c), which is unsafe.

- The banker's algorithm **considers each request as it occurs**, seeing whether granting it leads to a safe state. If it does, the request is granted; otherwise, it is postponed until later.



Banker's Algorithm (BA)


- When a new process enters the system, it must **declare the maximum number of instances of each resource type that it may need.**
- This number may **not exceed the total number of resources in the system.**
- When a user requests a set of resources, the system must determine whether the allocation of these resources will leave the system **in a safe state.**
- If it will, the resources are **allocated**; otherwise, the process must wait until some other process **releases enough resources**



Deadlock Avoidance

Banker's Algorithm (BA): Safe state, Basic facts

- Safe State: A system state where it's possible to find **a safe sequence for process execution that guarantees no deadlock.**
- If a system is in safe state \Rightarrow no deadlocks
- If a system is in unsafe state \Rightarrow possibility of deadlock
- Avoidance \Rightarrow ensure that a system will never enter an unsafe state.



Banker Algorithm Safe State

- Consider the following example for a single resource, assuming that 10 entities exist ($E=10$) for the given resource (Has (allocation), Max,):

Allocation

	Has	Max
A	3	9
B	2	4
C	2	7

Available

Free: 3
(a)

	Has	Max
A	3	9
B	4	4
C	2	7

Free: 1
(b)

	Has	Max
A	3	9
B	0	—
C	2	7

Free: 5
(c)

	Has	Max
A	3	9
B	0	—
C	7	7

Free: 0
(d)

	Has	Max
A	3	9
B	0	—
C	0	—

Free: 7
(e)

Fig: Demonstration that the state in (a) is safe (Tanenbaum)



Banker Algorithm Safe State

- Consider the following example (A gets an **extra resource** in (b)), again assuming that 10 instances exist for the given resource:
 - The state in *b* is **not yet deadlocked**, but **will deadlock eventually**.

Has Max		
A	3	9
B	2	4
C	2	7

Free: 3

(a)

Has Max		
A	4	9
B	2	4
C	2	7

Free: 2

(b)

Has Max		
A	4	9
B	4	4
C	2	7

Free: 0

(c)

Has Max		
A	4	9
B	—	—
C	2	7

Free: 4

(d)

Fig: Demonstration that the state in (b) is not safe. (Tanenbaum)



Banker's Algorithm (BA) Data Structure to implement BA

- We need the following data structures, where n is the number of processes in the system and m is the number of resource types:
 - **Allocation.** An $n \times m$ matrix defines the number of resources of each type currently **allocated to each process**. If $\text{Allocation}[i][j]$ equals k , then process P_i is currently allocated k instances of resource type R_j .
 - **Available.** A **vector of length m indicates the number of available resources** of each type. If $\text{Available}[j]$ equals k , then k instances of resource type R_j are available.
 - **Max.** An $n \times m$ **matrix defines the maximum demand of each process**. If $\text{Max}[i][j]$ equals k , then process P_i may request at most k instances of resource type R_j .
 - **Need.** An $n \times m$ matrix indicates **the remaining resource need of each process**. If $\text{Need}[i][j]$ equals k , then process P_i may need k more instances of resource type R_j to complete its task. Note that $\text{Need}[i][j]$ equals $\text{Max}[i][j] - \text{Allocation}[i][j]$

- Let say a system
 - has $n=2$ processes (p_0, p_1)
 - $m=2$ resources (printer and scanner), printer has 5 instance and scanner has 4 instance.
 - $E=[5,4]$ Available= $[2,1]$

	Allocation		Max	
	Pri	Sca	Pri	Sca
P_0	3	2	4	3
P_1	0	1	1	4

	Need	
	Pri	Sca
P_0	1	1
P_1	1	3



Deadlock avoidance

Banker's Algorithm: Steps in the Banker's Algorithm

1. Initialization:

1. Each process declares its maximum needs for each resource.
2. The system knows the total resources available and calculates which resources are allocated and which remain available.

2. Resource Request Handling:

1. When a process requests resources, the system checks:
 - 1. Request \leq Need:** The requested resources must be within the declared maximum needs.
 - 2. Request \leq Available :** The resources must be available.
2. If these checks pass, **the system temporarily allocates the resources**

3. Safety Check:

1. The system verifies if, after the allocation, there exists a safe sequence (a possible order for all processes to complete).
2. If a safe sequence exists, the allocation is confirmed. **If not, the process must wait.**



Baker Algorithm

Resource request Algorithm

- Determine whether requests can be safely granted.
- Let $Request_i$ be the request vector for process P_i . If $Request_i[j] == k$, then process P_i wants k instances of resource type R_j . When a request for resources is made by process P_i , the following actions are taken:

- 1) If $Request_i \leq Need_i$
Goto step (2) ; otherwise, raise an error condition, since the process has exceeded its maximum claim.
 - 2) If $Request_i \leq Available$
Goto step (3); otherwise, P_i must wait, since the resources are not available.
 - 3) Have the system pretend to have allocated the requested resources to process P_i by modifying the state as follows:
 - $Available = Available - Request_i$
 - $Allocation_i = Allocation_i + Request_i$
 - $Need_i = Need_i - Request_i$
- If safe : the resources are allocated to P_i
 - If unsafe: P_i must wait, and the old resource-allocation state is restored



Banker Algorithm Safety Algorithm

- A **safe state** is one where there is a feasible order for the processes to **finish without deadlock**.

```
1) Let Work and Finish be vectors of length 'm' and 'n' respectively.  
   Initialize: Work = Available  
   Finish[i] = false; for i=1, 2, 3, 4...n  
2) Find an i such that both  
   a) Finish[i] = false  
   b)  $Need_i \leq Work$   
   if no such i exists goto step (4)  
3) Work = Work + Allocation[i]  
   Finish[i] = true  
   goto step (2)  
4) if Finish [i] = true for all i then the system is in a safe state
```

- Consider a system with five processes **P0 through P4** and three resource types **A, B, and C**.
- Resource type A **has ten instances**, resource type B **has five instances**, and resource type C **has seven instances**.
- Suppose that, at time T0, the following snapshot of the system has been taken.
 - What will be the content of the **Need** matrix?
 - Is this state safe?
 - What will happen if process P1 requests one additional instance of resource type A and two instances of resource type C?

	Allocation			Max			Need			Available		
										3	3	2
P1	0	1	0	7	5	3						
P2	2	0	0	3	2	2						
P3	3	0	2	9	0	2						
P4	2	1	1	2	2	2						
P5	0	0	2	4	3	3						



Banker Algorithm

An Illustrative Example: Need Matrix

$\text{Need}[i, j] = \text{Max}[i, j] - \text{Allocation}[i, j]$ So, the content of Need Matrix is:

	Allocation			Max			Need			Available		
										3	3	2
P0	0	1	0	7	5	3	7	4	3			
P1	2	0	0	3	2	2	1	2	2			
P2	3	0	2	9	0	2	6	0	0			
P3	2	1	1	2	2	2	0	1	1			
P4	0	0	2	4	3	3	4	3	1			



Banker Algorithm

An Illustrative Example: Is it safe?

$m=3, n=5$ Step 1 of Safety Algo

Work = Available

Work =

3	3	2
---	---	---

0 1 2 3 4

Finish =

false	false	false	false	false
-------	-------	-------	-------	-------

For $i=0$ Step 2

Need₀ = 7, 4, 3

Finish [0] is false and Need₀ > Work

So P₀ must wait But Need ≤ Work

For $i=1$ Step 2

Need₁ = 1, 2, 2

Finish [1] is false and Need₁ ≤ Work

So P₁ must be kept in safe sequence

3, 3, 2 2, 0, 0 Step 3

Work = Work + Allocation₁

Work =

5	3	2
---	---	---

0 1 2 3 4

Finish =

false	true	false	false	false
-------	------	-------	-------	-------

For $i=2$ Step 2

Need₂ = 6, 0, 0

Finish [2] is false and Need₂ > Work

So P₂ must wait

For $i=3$ Step 2

Need₃ = 0, 1, 1

Finish [3] = false and Need₃ < Work

So P₃ must be kept in safe sequence

5, 3, 2 2, 1, 1 Step 3

Work = Work + Allocation₃

Work =

7	4	3
---	---	---

0 1 2 3 4

Finish =

false	true	false	true	false
-------	------	-------	------	-------

For $i=4$ Step 2

Need₄ = 4, 3, 1

Finish [4] = false and Need₄ ≤ Work

So P₄ must be kept in safe sequence

7, 4, 3 0, 0, 2 Step 3

Work = Work + Allocation₄

Work =

7	4	5
---	---	---

0 1 2 3 4

Finish =

false	true	false	true	true
-------	------	-------	------	------

For $i=0$ Step 2

Need₀ = 7, 4, 3

Finish [0] is false and Need ≤ Work

So P₀ must be kept in safe sequence

7, 4, 5 0, 1, 0 Step 3

Work = Work + Allocation₀

Work =

7	5	5
---	---	---

0 1 2 3 4

Finish =

true	true	false	true	true
------	------	-------	------	------

For $i=2$ Step 2

Need₂ = 6, 0, 0

Finish [2] is false and Need₂ ≤ Work

So P₂ must be kept in safe sequence

7, 5, 5 3, 0, 2 Step 3

Work = Work + Allocation₂

Work =

10	5	7
----	---	---

0 1 2 3 4

Finish =

true	true	true	true	true
------	------	------	------	------

Finish [i] = true for $0 \leq i \leq n$ Step 4

Hence the system is in Safe state

The safe sequence is P₁, P₃, P₄, P₀, P₂

	Allocation			Max			Need			Available		
										3	3	2
P0	0	1	0	7	5	3	7	4	3			
P1	2	0	0	3	2	2	1	2	2			
P2	3	0	2	9	0	2	6	0	0			
P3	2	1	1	2	2	2	0	1	1			
P4	0	0	2	4	3	3	4	3	1			

Banker Algorithm

An Illustrative Example: What will happen if process P1

- What will happen if process P1 requests one additional instance of resource type A and two instances of resource type C?

	Allocation			Max			Need			Available		
										2	3	0
P0	0	1	0	7	5	3	7	4	3			
P1	3	0	2	3	2	2	0	2	0			
P2	3	0	2	9	0	2	6	0	0			
P3	2	1	1	2	2	2	0	1	1			
P4	0	0	2	4	3	3	4	3	1			

Request₁ = $\begin{matrix} A & B & C \\ 1 & 0 & 2 \end{matrix}$

To decide whether the request is granted we use Resource Request algorithm

Step 1
 $\begin{matrix} 1, 0, 2 & 1, 2, 2 \\ \text{Request}_1 \leq \text{Need}_1 \end{matrix}$ ✓

Step 2
 $\begin{matrix} 1, 0, 2 & 3, 3, 2 \\ \text{Request}_1 \leq \text{Available} \end{matrix}$ ✓

Step 3

Available = Available – Request₁
Allocation₁ = Allocation₁ + Request₁
Need₁ = Need₁ - Request₁

Process	Allocation	Need	Available
	A B C	A B C	A B C
P ₀	0 1 0	7 4 3	2 3 0
P ₁	3 0 2	0 2 0	
P ₂	3 0 2	6 0 0	
P ₃	2 1 1	0 1 1	
P ₄	0 0 2	4 3 1	

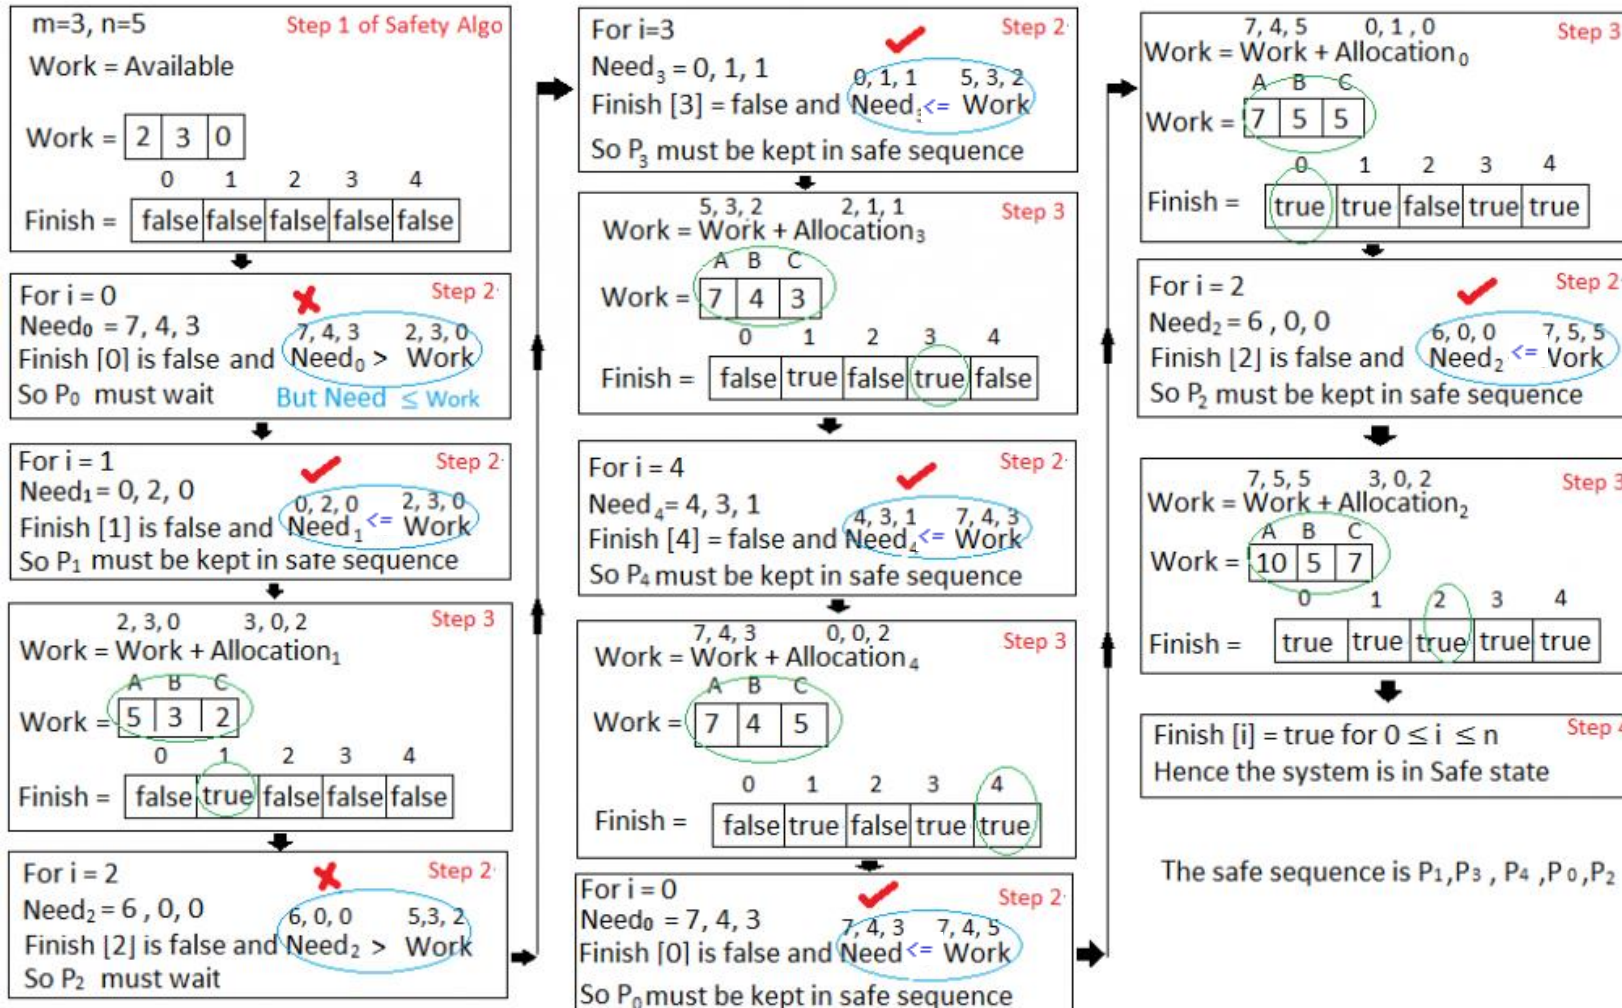
$[3, 3, 2] - [1, 0, 2] = [2, 3, 0]$
 $[2, 0, 0] + [1, 0, 2] = [3, 0, 2]$
 $[1, 2, 2] - [1, 0, 2] = [0, 2, 0]$



Banker Algorithm

An Illustrative Example: What will happen if process P1

- We must determine whether this new system state is safe. To do so, we again execute **Safety algorithm** on the above data structures.



	Allocation			Max			Need			Available		
										2	3	0
P0	0	1	0	7	5	3	7	4	3			
P1	3	0	2	3	2	2	0	2	0			
P2	3	0	2	9	0	2	6	0	0			
P3	2	1	1	2	2	2	0	1	1			
P4	0	0	2	4	3	3	4	3	1			



Homework

- When the system is in this state if P4 request for R(3,3,0). Can it be granted?

	Allocation			Max			Need			Available		
										2	3	0
P0	0	1	0	7	5	3	7	4	3			
P1	3	0	2	3	2	2	0	2	0			
P2	3	0	2	9	0	2	6	0	0			
P3	2	1	1	2	2	2	0	1	1			
P4	0	0	2	4	3	3	4	3	1			

- P4 cannot be granted, since the resources are not available.



Homework

- A request for (0,2,0) by P0 cannot be granted, even though the resources are available, since the resulting state is unsafe. **Show it. Home Work!**

	Allocation			Max			Need			Available		
										2	3	0
P0	0	1	0	7	5	3	7	4	3			
P1	3	0	2	3	2	2	0	2	0			
P2	3	0	2	9	0	2	6	0	0			
P3	2	1	1	2	2	2	0	1	1			
P4	0	0	2	4	3	3	4	3	1			

- Modern Operating Systems (Tanenbaum): **Chapter 6(6.5)**
- Operating System Concepts (Silberschatz): **Chapter 7(7.5)**
- Operating Systems: Internals and Design Principles (Starlings): **Chapter 6(6.3)**