

What is File System Management

- Goals for Today Overview
- User view of file systems
- System calls
- Structures, organisation, file types
- Implementation view of file systems
- Disk and partition layout
- File tables
- Free space management ...

Different views of file system

- File Systems Different Views
- A user view that defines a file system in terms of the abstractions that the operating system provides
 - How files are named, what operations allowed on them
 - What directories look like
 - Interface issue
- An implementation view that defines the file system in terms of its low level implementation
 - How files and directories are stored
 - How disk space is managed
 - How to make everything work efficiently



Figure: User vs. Implementation View

Implementation view=>context of file system

- Implementation view of file system Context
- Irrespective of the type of file system, a number of additional considerations have to be addressed, including
 - Disk partitions, partition tables, boot sectors, etc.
 - Free space management (= free memory)
 - System wide and per process file tables (= process tables)
 - How file and directories store (to be introduced in the next lecture)
- Low level formatting writes sectors to the disk, high level formatting imposes a file system on top of this
 - File system manage blocks that can cover multiple sectors
 - The file is split up into a number of (not necessarily) contiguous blocks.

File types

- Files Types
- Many OSs support several types of file.
- Both Windows and Unix (including OS X) have regular files and directories.
- Regular files contain user data in ASCII or binary (well defined) format
 - A file is a named collection of data
 - ASCII files consist of lines of text, (can be displayed, can be edited by text editor)
 - Binary files have internal structure
- Directories group files together (but are files on an implementation level)
 - They provide a mapping of the logical file into the physical location

Directories

- Directories Directory structure
- Directories are special files that group files together and of which the structure is defined by the file system
 - A bit is set to indicate that they are directories
- Different directory structures have been used over the years
 - Single level: all files in the same directory (reborn in consumer electronics)
 - Two or multiple level directories (hierarchical) tree structures
 - Absolute path name from the root of the file system
 - Relative path name: the current working directory is used as the starting point
 - Directed acyclic graph (DAG): allows files to be shared (i.e. links to files or sub-directories) but cycles are forbidden

Implementation => How Disk is structured

- Hard Disk Structures Partitions
- Disks are usually divided into multiple partitions
 - An independent file system may exist on each partition
- Master boot record located at start of the entire drive
 - Used to boot the computer (BIOS reads and executes MBR)
 - Contains partition table as it is and with active partition
 - One partition is listed as active containing a boot block to load the operating system

Implementation => How free disk space is managed?

- Disk Space Management Free Space Management
- Two methods are commonly used to keep track of free disk space: **Bitmaps** and **Linked Lists**
 - Note that these approaches are very similar to the ones to keep track of free memory

Implementation => How files are managed?

- File managed by OS Implementation
- Apart from the free disk space tables, there is a number of key data structures stored in memory:
 - An in-memory mount table
 - An in-memory directory cache of recently accessed directory information
 - A system-wide open file table, containing a copy of the FCB for every currently open file in the system, including location on disk, file size, and "open count" (processes that use the file)
 - File control blocks (FCBs) are kernel data structures, i.e. they are protected and only accessible in kernel space
 - A per-process open file table, containing a pointer to the system open file table

File/ Directory Implementation => How files and directories are stored? (next lecture)

Implementation => How shared files are managed? (next lecture)

File system calls

- System Calls
- System calls enable a user application to ask the operating system to carry out an action on its behalf (in kernel mode)
- There are two different categories of system calls:
 - File manipulation: open(), close(), read(), write(), delete(), ...
 - Directory manipulation: create(), delete(), read(), rename(), link(), unlink(), list(), update()

Directory structure

- Directories Directory structure
- Figure: Single-level directory and DAG Directory Implementation (Tavernbaum)

- Directories DAG and Graph Complications
- The use of DAGs results in significant complications in the implementation
- Files have multiple absolute file names
- Deleting files becomes a lot more complicated (i.e. links may no longer point to a file)
- A garbage collection scheme may be required to remove files that are no longer accessible from the file system tree

Partition layout

- Partition Layouts A Unix Partition
- The layout of a partition differs depending on the file system
 - The partition boot block:
 - Contains code to boot the operating system
 - Every partition has boot block - even if it does not contain OS
 - Super block contains the partition's file system's key parameters, e.g., partition size, number of blocks, inode table size
 - Read into main memory when computer is booted
 - Free space management contains, e.g., a bitmap or linked list that indicates the free blocks
 - Inodes contains information on files, commonly maintained in inodes
 - Root directory contains the top of the file-system tree
 - Data: files and directories

Bitmap vs. Linked list

- Disk Space Management Free Space Management
- Bitmaps represent each block by a single bit in a map
 - The size of the bitmap grows with the size of the disk but is constant for a given disk
 - Bitmaps take comparably less space than linked lists
 - 1 bit per block vs. 1 pointer (32/64/48) per block
- Linked Lists:
 - Free blocks are used to hold the numbers of the free blocks (hence, they are no longer free)
 - Free blocks are linked when the disk becomes full, this is not wasted space
 - Blocks are linked together, i.e., multiple blocks list the free blocks
 - The size of the list grows with the size of the disk and shrinks with the size of the blocks (not constant)
 - Linked lists can be modified by keeping track of the number of consecutive free blocks for each entry

File table management

- File managed by OS File Tables
- Figure: File control block (FCB) and File table

Directories system calls

- Directories System Calls
- Similar to files, directories are manipulated using system calls
 - create/delete: a new directory is created/deleted
 - opendir, closedir: add/free directory to/from internal tables
 - Others: rename, link, unlink, list, update
- libfsck (fsck) means copy the file and its directory to its own directory using the command

Example

- Disk Space Management Free Space Management: bitmap vs. linked list
- Bitmaps:
 - Require extra space. E.g., if block size = 2¹² bytes and disk size = 2²⁸ bytes (1 GB) => bitmap size = 2¹⁶ (256K)
 - Proportional to number of block
 - Keeping it in main memory is possible only for small disks.
- Linked lists:
 - No waste of disk space
 - The only need to keep in memory one block of pointers (load a new block when need)
 - Block size = 2¹² (4K) bytes, 32-bit disk block number and disk size = 2²⁸ bytes
 - Block number: 0 <= i <= 2¹⁶ - 1
 - linked list block number = 2¹² * i + 1
 - linked list size = 2¹⁶ * 488 = 152448

File table usage

- File Tables Illustration
- Figure: (a) Opening a file (b) reading a file (c) releasing a file (d) releasing a file