



Operating Systems and Concurrency

Lecture 23:
File System 3

University of Nottingham, Ningbo China
2024



Goals for Today

Overview

- **File implementations**
 - Contiguous
 - Linked lists
 - File Allocation Table (FAT)
 - I-nodes
- **Directories implementation**
 - I-nodes lookup
- **Hard and soft links**

File access

Sequential vs. Random Access

- Files will be composed of a number of blocks.
- Files are **sequential** or **random access**. Random access is essential for example in **database systems**.

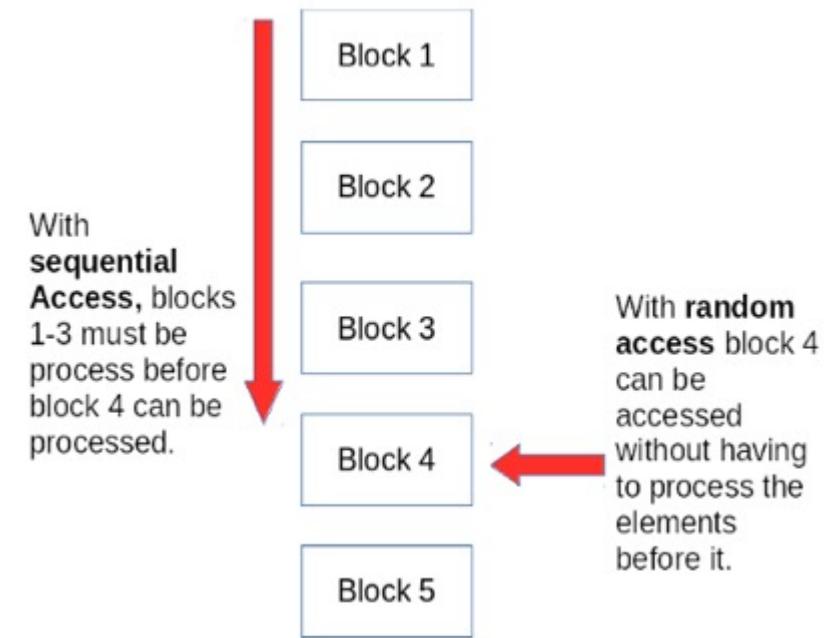


Figure: Types of access to a file

Contiguous Allocation

Concept

- **Contiguous file systems** are similar to **dynamic partitioning** in memory allocation:
- Each file is stored in a single group **of adjacent blocks** on the hard disk E.g. 1KB blocks, 100KB file, we need 100 contiguous blocks
- Allocation of free space can be done using **first fit, best fit, next fit**, etc.

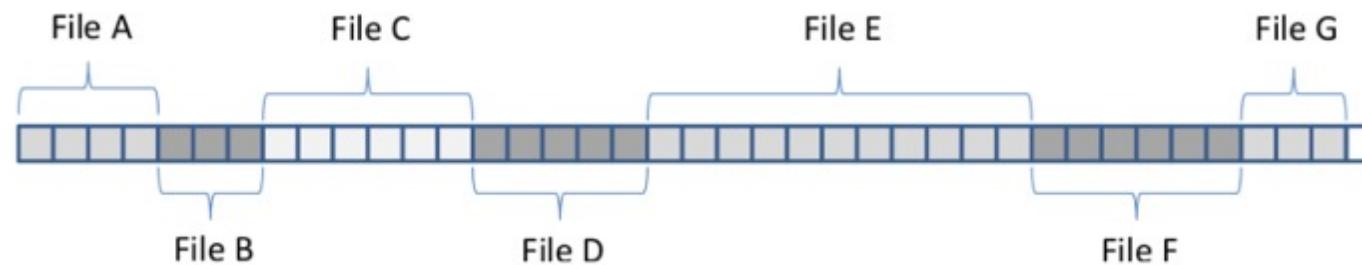


Figure: External Fragmentation when Removing Files

Contiguous Allocation

Advantages

- **Simple to implement:** only location of the **first block** and the **length** of the file must be stored (in the directory entry)
- **Optimal read/write performance:** blocks are co-located/clustered in nearby/adjacent sectors, hence the seek time is minimised (remember the example in lecture on disks!)

File	Start	Length
req1.c	0	12
req1.o	30	10
req1	15	5
req1.txt	41	20

Figure: Directory table



Contiguous Allocation

Disadvantages

- Disadvantages of contiguous file systems include:
 - The **exact size** of a file (process) is not always **known beforehand**: what if the file size exceeds the initially allocated disk space
 - **Allocation algorithms** needed to decide which free blocks to allocate to a given file (e.g., first fit, best fit, etc.)
 - Deleting a file results in **external fragmentation**: de-fragmentation must be carried out regularly (and is slower than for memory)
 - If a single block in a contiguous file becomes **corrupted**, the entire file can become **unreadable**, as all data is stored in a continuous block.
- Contiguous allocation is still in use: **CD-ROMS/DVDs**
 - External fragmentation is less of an issue here since they are write once only

Linked Lists

Concept

- To **avoid external fragmentation**, files are stored in **separate blocks** (similar to paging) that are **linked to one another**
- Only the **address of the first** block has to be stored to locate a file
- Each block contains a **data pointer** to the next block (which takes up space)

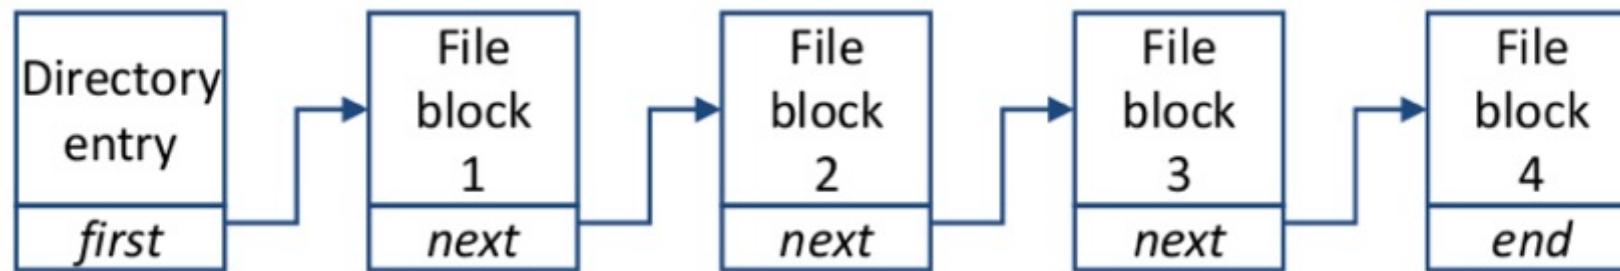


Figure: Linked List File Storage



Linked Lists

Advantages

- **Easy to maintain:** only the first block (address) has to be maintained in the **directory entry**
- File sizes can **grow dynamically** (i.e. file size does not have to be known beforehand): new blocks/sectors can be added to the end of the file
- Similar to paging for memory, every possible block/sector of disk space can be used: i.e., there is **no external fragmentation!**
- **Sequential access** is straightforward, although **more seek operations/disk access** may be required



Linked Lists

Disadvantages

- **Random access is very slow**, to retrieve a block in the middle, one has to walk through the list from the start
- Space is lost within the blocks due to the pointer, the data in a **block is no longer a power of 2!**
- **Diminished reliability:** if one block is corrupt/lost, access to the rest of the file is lost

File Allocation Tables

Key Concept

- Store the linked-list pointers in a **separate index table**, called a **File Allocation Table (FAT)**, in memory!

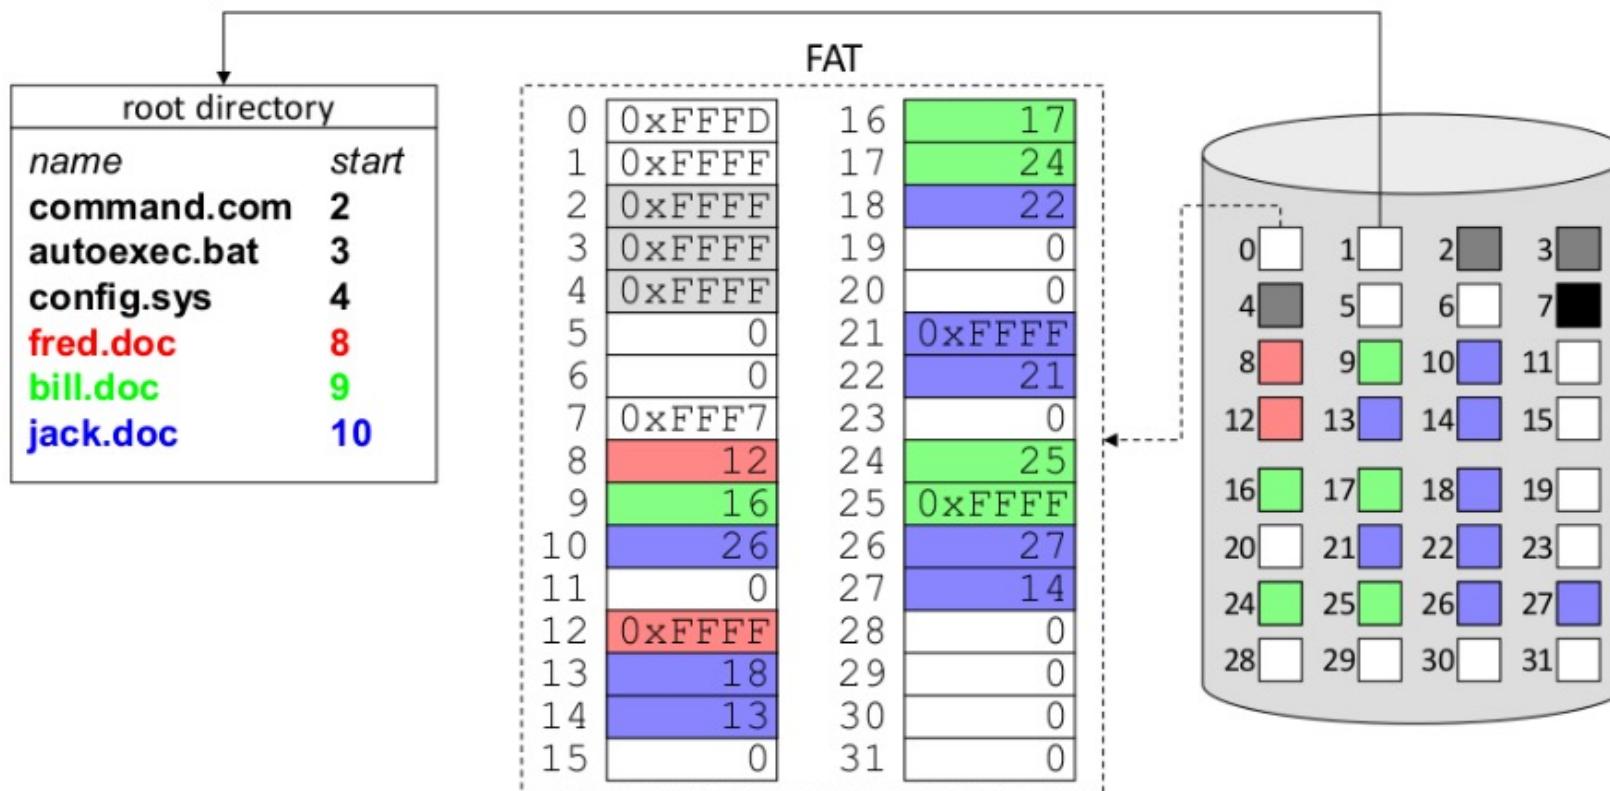


Figure: File Allocation Tables



File Allocation Tables

Advantages and disadvantages

- Advantages:
 - **Block size remains power of 2**, i.e., no space is lost due to the pointer
 - **Index table** can be **kept in memory** allowing fast **non-sequential/random** access (one still has to walk through the table though)
- Disadvantages:
 - The **size of the file allocation table grows** with the number of blocks, and hence the size of the disk
 - For a 200GB disk, with a 1KB block size, 200 million entries are required. Assuming that each entry at the table occupies 4 bytes, this requires **800Mb of main memory!**



- **Each file** has a small data structure (on disk) called **I-node** (index-node) that contains its **attributes** and **block pointers**.
 - In contrast to FAT, an I-node is **only loaded when the file is open** (stored in system wide open file table)
 - If every I-node consists of n bytes, and at most k files can be open at any point in time, at most $n \times k$ bytes of main memory are required
- I-nodes are composed of **direct block pointers** (usually 10), **indirect block pointers**, or a combination thereof (e.g., similar to **multi-level page tables**)

I-nodes

Concept

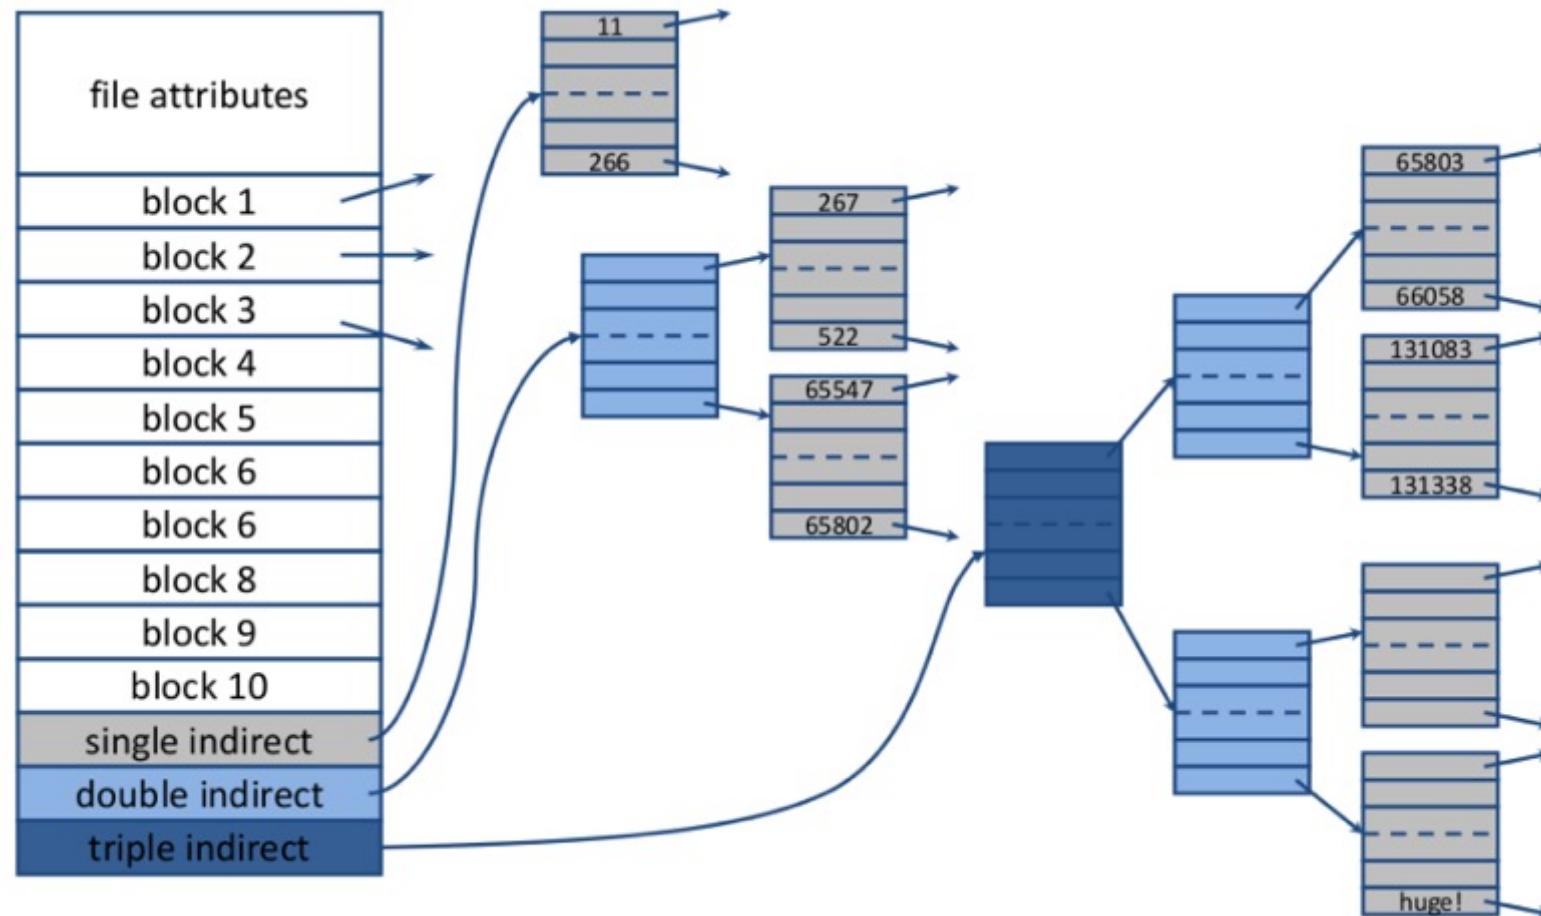


Figure: File Storage Using I-Nodes



- Assuming a block size of **1KB**, and **32-bits disk address space**. With only direct block pointers, the maximum size of file could be $1\text{KB} \times \text{number of direct blocks}$ (e.g 10, a total of 10KB).
- A single indirect block can store up to **256 block pointers** ($32 \text{ bits} = 4 \text{ bytes per pointer}; 2^{10}/2^2 = 2^8 = 256$ pointers). That is, with 10 direct blocks + 1 indirect block, we can have files with up to 266 blocks ($\Rightarrow 266\text{KBs}$).
- A double indirect points to a block of 256 block pointers. Each of which points to 256 indirect blocks. Therefore, with the additional one double indirect block, we could have files with size up to $266\text{KBs} + (256 \times 256 = 65536) \Rightarrow 65802 \text{ KBs}$.
- If we need files larger than 64M, we will need **a triple indirect**.



Directories

Possible Implementations

- Directories contain a list of **human readable file names** that are mapped onto **unique identifiers** and **disk locations**
 - They provide a **mapping** of the **logical file** onto the **physical location**
- Retrieving a file comes down to **searching a directory file** as fast as possible:
 - A simple **random order** of **directory entries** might be insufficient (search time is linear as a function of the number of entries)
 - Indexes or **hash tables** can be used

Directories

Possible Implementations

- They can store all file related **attributes** (e.g. file name, disk address – Windows) or they can contain **a pointer** to the data structure (i-node) that contains the details of the file (Unix)

File Name	Attributes
...	...
...	...
...	...
...	...
...	...
...	...
...	...
...	...
...	...

File Name	Pointer
...	...
...	...
...	...
...	...
...	...
...	...
...	...
...	...
...	...

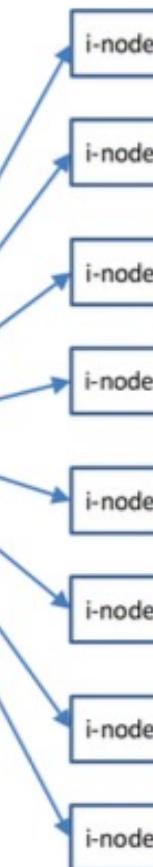


Figure: Directory Implementations

Directories

Implementation with i-nodes

- In **UNIX**, all information about the file (type, size, date, owner, and block pointers) is stored in its i-node.
- Therefore, **directory tables** are very simple data structures composed of file name and a pointer to the i-node.
- Note that directories are no more than **a special kind of file**, so they have their **own i-node**.

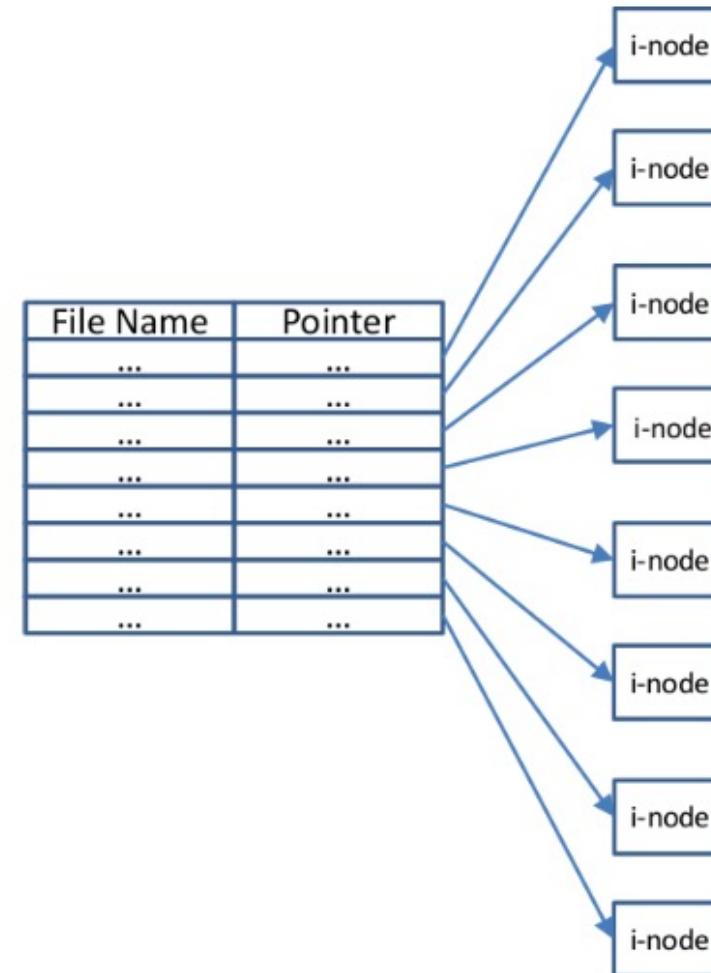


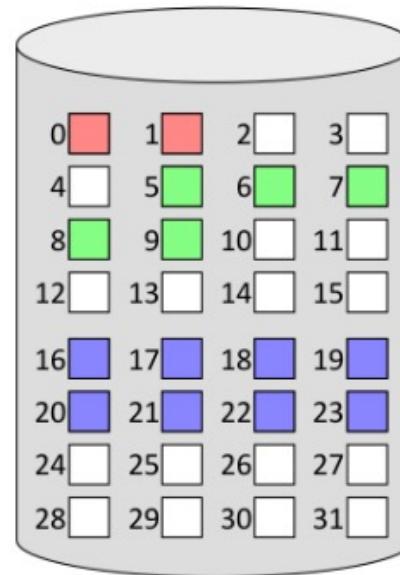
Figure: I-node Directory Structure

File System Comparison

Contiguous vs. Linked vs. Indexed

contiguous

directory		
<i>name</i>	<i>start</i>	<i>size</i>
fred.doc	0	2
bill.doc	5	5
jack.doc	16	8



linked list

directory	
<i>name</i>	<i>start</i>
fred.doc	0
bill.doc	1
jack.doc	2



indexed

directory	
<i>name</i>	<i>index</i>
fred.doc	0
bill.doc	8
jack.doc	16

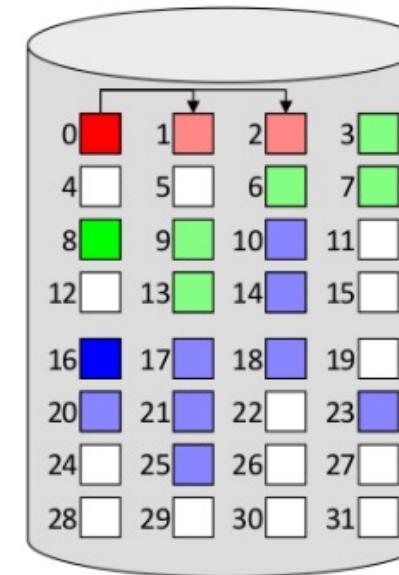


Figure: Contiguous vs. Linked List (or FAT) vs. I-nodes (or indexed)

I-nodes

Lookups

- **Opening a file** requires the **disk blocks** to be **located**
 - **Absolute file names** are located relative to the root directory
 - **Relative file names** are located based on the current working directory
- E.g. Try to locate /usr/gdm/mbox

/ inode	/ contents	/usr inode	/usr contents	/usr/gdm inode	/usr/gdm contents
1	2	6	132	26	406
size	1 .	size	6 .	size	26 .
mode	1 ..	mode	1 ..	mode	6 ..
times	4 bin	times	19 fred	times	64 research
2	7 dev	132	30 bill	406	92 teaching
	14 lib		51 jack		60 mbox
	9 etc		26 gdm		17 grants
	6 usr		45 cfi		
	8 tmp				

Figure: Locating a File



Locate the root directory of the file system

- Its i-node sits on a fixed location at the disk (the directory itself can sit anywhere)

Locate the directory entries specified in the path:

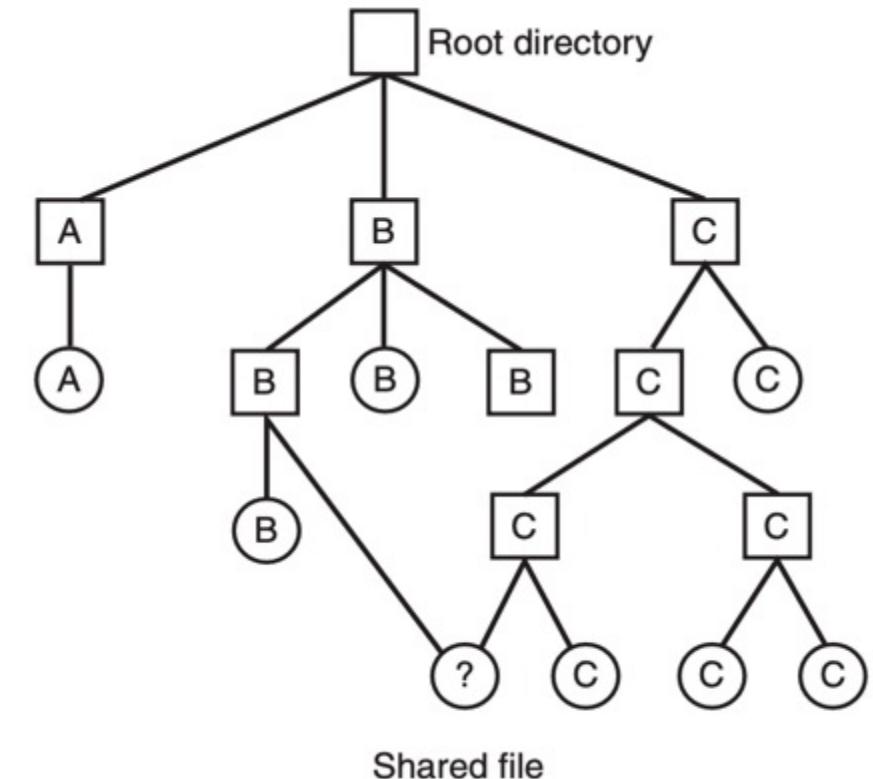
- Locate the i-node number for the first component (directory) of the path that is provided
- Use the i-node number to index the i-node table and retrieve the directory file
- Look up the remaining path directories by repeating the two steps above

Once the file's directories have been located, locate the file's i-node and cache it into memory

I-nodes: Sharing files between directories

Hard and Soft Links

- In **Directed acyclic graph (DAG)**, there are two approaches to **share a file**, e.g. between directory B and C, where C is the 'real' owner:
 - **Hard links:**
 - maintain two (or multiple) references to the same i-node in B and C.
(In Unix: ln /C/file1 /B/file2)
 - the i-node link reference counter will be set to 2
 - **Symbolic(soft) links:**
 - The owner maintains a reference to the i-node in, e.g., directory C
 - The "referencer" maintains a small file (that has its own i-node) that contains the location and name of the shared file in directory C.
(In Unix: ln -s /C/file1 /B/file2)
- What is the **best approach?** ⇒ both have advantages and disadvantages



I-nodes: Sharing files between directories

Hard Links

- **Hard links** are the **fastest way of linking files!**
- Disadvantages of hard links:
 - Assume that the **owner** of the file **deletes** it:
 - If the i-node is also **deleted**, any hard link will, in the best case, **point to an invalid i-node**
 - If the i-node gets **deleted** and “**recycled**” to point to **another file**, the hard links will **point to the wrong file!**
 - The only solution is to **delete the file**, and **leave the i-node intact** if the “**reference count**” is larger than 0 (the original owner of the file still gets “charged” for the space until B decides to remove it)

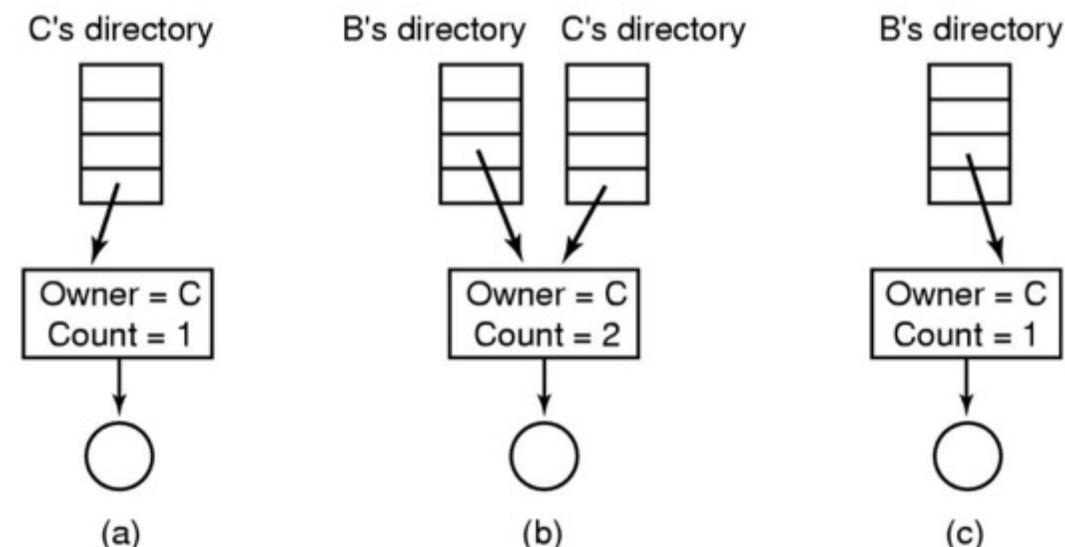


Figure: (a) Before creating a link. (b) After creating a link. (c) after the original owner removes the file



- Disadvantages of **soft links**:
 - They result in an **extra file lookup** (once the link file has been found, the original file needs to be found as well)
 - They require an **extra i-node** for the link file
- Advantages of soft links:
 - There are **no problems with deleting** the original file ⇒ the file simply does not exist any more
 - They can **cross the boundaries of machines**, i.e. the linked file can be located on a different machine



I-nodes

Hard and soft Links in Unix

```
Desktop Documents Downloads Music Pictures Public Templates Videos
parallels@ubuntu-linux-22-04-desktop:~$ cd Documents/
parallels@ubuntu-linux-22-04-desktop:~/Documents$ ls
test1.c
parallels@ubuntu-linux-22-04-desktop:~/Documents$ ln ./test1.c hltest
parallels@ubuntu-linux-22-04-desktop:~/Documents$ ln -s ./test1.c sltest
parallels@ubuntu-linux-22-04-desktop:~/Documents$ ls
hltest sltest test1.c
parallels@ubuntu-linux-22-04-desktop:~/Documents$ ls -ali
total 16
3015658 drwxr-xr-x  2 parallels parallels 4096 Nov 20 19:53 .
2621442 drwxr-x--- 16 parallels parallels 4096 Nov 20 19:52 ..
3014690 -rw-rw-r--  2 parallels parallels     4 Nov 20 19:52 hltest
3014686 lrwxrwxrwx  1 parallels parallels     9 Nov 20 19:53 sltest -> ./test1.c
3014690 -rw-rw-r--  2 parallels parallels     4 Nov 20 19:52 test1.c
parallels@ubuntu-linux-22-04-desktop:~/Documents$ █
```

What will happen if I delete test1.c and open the two links?



Summary

Take-Home Message

- Contiguous, linked list, FAT and i-nodes as file system implementations.
- Lookups with I-nodes.
- Hard and Soft Links