


Presentation for use with the textbook **Data Structures and Algorithms in Java, 6th edition**, by M. T. Goodrich, R. Tamassia, and M. H. Goldwasser, Wiley, 2014

Binary Search Trees



Reading

M. T. Goodrich, R. Tamassia and M. H. Goldwasser,
Data Structures and Algorithms in Java, 6th Edition,
2014.

- **Chapter 10. Hash Tables, Maps, and Skip Lists**
- **Section 10.3 The Sorted Map ADT**
- **pp. 396-401**

- **Chapter 11. Search Tree Structures**
- **Sections 11.1-11.2**
- **pp. 423-442**

Learning Objectives

- ◆ To be able to understand and describe the Sorted Map ADT;
- ◆ To be able to analyze the complexity of the Sorted Map ADT methods;
- ◆ To be able to implement the Sorted Map ADT with a binary search tree;
- ◆ To be able to explain the update operations for a binary search tree;
- ◆ To be able to apply the Sorted Map ADT and binary search tree.

Motivation

Suppose you have an array of integers and want to search for a particular integer k , how will you do it?

线性查找 (Linear Search) 对无序数组适用时间复杂度: $O(n)$


二分查找 (Binary Search) 要求数组是有序的时间复杂度: $O(\log n)$

使用哈希表 (Hash Table) 如果提前建立映射, 查找效率可以达到 $O(1)$

构建搜索树 (如 BST / AVL / 红黑树) 插入与查找都能达到 $O(\log n)$ (平衡状态下)

Motivation

- ◆ If the array is not sorted, then we need to scan all elements, hence $O(n)$.
- ◆ If the array is sorted, then we can do better. How?



Ordered Maps

- ◆ Keys are assumed to come from a total order.
- ◆ Items are stored in order by their keys
- ◆ This allows us to support nearest neighbor queries: 最近邻查询
 - ◆ Item with largest key less than or equal to k
 - ◆ Item with smallest key greater than or equal to k
 - 找到小于等于某个键 k 的最大键项
 - 找到大于等于某个键 k 的最小键项

The Sorted Map ADT (Sec.10.3)

The Sorted Map ADT includes all methods of the Map ADT, plus the following.

- ❑ **firstEntry()**: Returns the entry with **smallest key value** (or null, if the map is empty).
- ❑ **lastEntry()**: Returns the entry with **largest key value** (or null, if the map is empty).
ceilingEntry(k)返回最小的键 k 的项 (若不存在则返回 null)
例如: map = {3, 7, 10}, 查询 ceilingEntry(8) 10
- ❑ **ceilingEntry(k)**: Returns the entry **with the least key value greater than or equal to k** (or null, if no such entry exists).
- ❑ **floorEntry(k)**: Returns the entry **with the greatest key value less than or equal to k** (or null, if no such entry exists).
floorEntry(k)返回最大的键 k 的项 (若不存在则返回 null)
例如: map = {3, 7, 10}, 查询 floorEntry(8) 7

The Sorted Map ADT

返回严格小于 k 的最大键项（如果不存在，则返回 null）
例：若 map 为 {3, 7, 10}，查询 lowerEntry(7) 得到 3

- ❑ **lowerEntry(k)**: Returns the entry with the greatest key value strictly less than k (or null, if no such entry exists).
返回严格大于 k 的最小键项（如果不存在，则返回 null）
例：若 map 为 {3, 7, 10}，查询 higherEntry(7) 得到 10
- ❑ **higherEntry(k)**: Returns the entry with the least key value strictly greater than k (or null if no such entry exists).
- ❑ **subMap(k_1, k_2)**: Returns an iteration of all entries with key greater than or equal to k_1 , but strictly less than k_2 .


返回一个迭代器，遍历所有键值 k 且 $k_1 \leq k < k_2$ 的项即： $[k_1, k_2)$
例：map 为 {2, 4, 6, 8, 10}，查询 subMap(4, 9) 得到 {4, 6, 8}

Sorted Search Tables


- ◆ We store the map's entries in an array list *A* so that they are in increasing order of their keys.
- ◆ We refer to this implementation as a ***sorted search table***.

0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
2	4	5	7	8	9	12	14	17	19	22	25	27	28	33	37


Binary Search



- ◆ Binary search can perform nearest neighbor queries on an ordered map that is implemented with an array, sorted by key
 - similar to the high-low children's game
 - at each step, the number of candidate items is halved
 - terminates after $O(\log n)$ steps
- ◆ Example: **find(7)**



Sorted Search Tables



- ◆ A search table is an ordered map implemented by means of a sorted sequence
 - We store the items in an array-based sequence, sorted by key
 - We use an external comparator for the keys
- ◆ Performance:
 - Searches take $O(\log n)$ time, using binary search
 - Inserting a new item takes $O(n)$ time, since in the worst case we have to shift n items to make room for the new item
 - Removing an item takes $O(n)$ time, since in the worst case we have to shift n items to compact the items after the removal
- ◆ The lookup table is effective only for ordered maps of *small size* or for maps on which *searches* are the most common operations, while insertions and removals are rarely performed (e.g., credit card authorizations)

Performance of a sorted map

(implemented using a sorted search table)

Method	Running Time
size	$O(1)$
get	$O(\log n)$
put	$O(n)$; $O(\log n)$ if map has entry with given key
remove	$O(n)$
firstEntry, lastEntry	$O(1)$
ceilingEntry, floorEntry, lowerEntry, higherEntry	$O(\log n)$
subMap	$O(s + \log n)$ where s items are reported
entrySet, keySet, values	$O(n)$

Motivation

- ◆ Binary search on ordered arrays is efficient: $O(\log_2 n)$
- ◆ However, insertion or removal of an item in an ordered array is slow: $O(n)$
- ◆ Ordered arrays are best suited for *static* searching, where search space does not change.
- ◆ Binary search trees can be used for *efficient dynamic searching*.

Binary Search Trees

A binary search tree is a *proper* binary tree storing keys (or key-value entries) at its *internal nodes* and satisfying the following property:

- Let u , v , and w be three nodes such that u is in the left *subtree* of v and w is in the right *subtree* of v . We have $key(u) \leq key(v) \leq key(w)$
- Assuming there are no duplicate keys, we have $key(u) < key(v) < key(w)$

• 对于任意节点 v :

• 所有在其左子树的节点 u , 满足:


$key(u) \leq key(v)$

• 所有在其右子树的节点 w , 满足:

$key(v) \leq key(w)$

• 如果不允许重复键值 (常见设定), 则满足严格不等:


$key(u) < key(v) < key(w)$



Binary Search Trees


External nodes do not store items

- and likely are not actually implemented, but are just null links from the parent




Is this a binary search tree?

一个二叉搜索树需要满足：
对于任意节点 v ：
所有左子树节点的值 $< v$
所有右子树节点的值 $> v$




Is this a binary search tree?

No



Is this a binary search tree?

No



从根节点开始，沿着树向下查找。

每一步：

如果 $k <$ 当前节点值 去左子树


如果 $k >$ 当前节点值 去右子树 如果 $k ==$ 当前节点值 找到目标

如果走到底仍未找到（即遇到 null 节点），则目标不存在。

Search

- ◆ To search for a key k , we trace a downward path starting at the root
- ◆ The next node visited depends on the comparison of k with the key of the current node
- ◆ If we reach a leaf, the key is not found
- ◆ Exercise: Write down the pseudo-code of the Search method.

Algorithm *Node TreeSearch(Key k , Node n)*



Search

Algorithm *Node TreeSearch*(Key k , Node n)

if $n.isExternal()$ // or “if $n == null$ ”

return *null*

if $k < n.key()$

return *TreeSearch*(k , $n.left()$)


else if $k = n.key()$

return n

else // $k > n.key()$

return *TreeSearch*(k , $n.right()$)


- ◆ Example: *get*(4):
 - Call *TreeSearch*(4, root)
- ◆ The algorithms for nearest neighbor queries are similar



Fundamental Property of Binary Search Trees


中序遍历的顺序是：左子树 当前节点 右子树
对每个节点递归执行此顺序。1, 2, 4, 6, 8, 9

- ◆ What is an inorder traversal of a tree?
- ◆ Exercise: what does an inorder traversal of the following search tree produce?



Fundamental Property of Binary Search Trees

An inorder traversal of a binary search tree visits the keys in increasing order.




Fundamental Property of Binary Search Trees

- ◆ How to access the minimal key? 从根开始，不断往左子节点走，直到没有左子节点。
- ◆ How to access the maximal key? 从根开始，不断往右子节点走，直到没有右子节点。
- ◆ How to access the largest key less than or equal to k ?
- ◆ How to access the smallest key greater than or equal to k ?

3. 从根向下搜索：若当前节点值 $> k$ ，往左走；
若当前节点值 $= k$ ，当前节点可能是候选解，还需去右子树看看有没有更大但仍 $\leq k$ 的值。

4. 从根向下搜索：若当前节点值 $< k$ ，往右走；
若当前节点值 $= k$ ，当前可能是候选解，继续向左走尝试找到更小的。



Insertion

- ◆ Exercise: How to do insertion, *put*(k, o)?
- ◆ Have to insert k where a *get*(k) would find it.
- ◆ So natural that *put*(k, o) starts with *get*(k)


从根节点开始；
如果 $k <$ 当前节点 去左子树；
如果 $k >$ 当前节点 去右子树；
如果遇到空位置（即 None），就在那里插入新的节点；
如果遇到相同的 key 选择更新值（可选逻辑）。

Insertion



- ◆ Exercise: How to do insertion, *put*(k, o)?
- ◆ We search for key k (using TreeSearch)
- ◆ Assume k is already in the tree then just replace the value.
- ◆ Otherwise, let w be the leaf reached by the search, we insert k at node w and expand w into an internal node

Insertion

◆ Example: insert 5




Insertion



Deletion

◆ How can we perform the operation *remove(k)*?


- E.g., remove 7, 5, 4, or 2



Deletion

上图（中间）：
4 是 2 的右子节点，5 是 4 的右子；
4 的左子节点为空（外部节点）；
所以可以将 4 和它的左“空子节点”一起删除；
然后将 5 提上，直接连接到 2 的右子上。



- ◆ To perform operation *remove*(k), we search for key k
- ◆ Assume key k is in the tree, and let v be the node storing k
- ◆ If node v has a leaf child w , we remove v and w from the tree with operation *removeExternal*(w), which removes w and its parent
- ◆ Example: remove 4



Deletion (cont.)

- ◆ We consider the case where the key k to be removed is stored at a node v whose children are both internal
 - we find the internal node w that follows v in an **inorder traversal**
 - we copy $key(w)$ into node v
 - we remove node w and its left child z (which must be a leaf) by means of operation *removeExternal*(z)

◆ Example: remove 3



Self-study: can we let w be the internal node that precedes v in an inorder traversal? Then how to perform "delete"? pp. 428-429

Performance

- ◆ Consider an ordered map with n items implemented by means of a binary search tree of height h

- the space used is $O(n)$
- methods **get**, **put** and **remove** take $O(h)$ time

- ◆ The height h is $O(n)$ in the worst case and $O(\log n)$ in the best case

