


Presentation for use with the textbook **Data Structures and Algorithms in Java, 6th edition**, by M. T. Goodrich, R. Tamassia, and M. H. Goldwasser, Wiley, 2014

Binary Search Trees



Reading

M. T. Goodrich, R. Tamassia and M. H. Goldwasser,
Data Structures and Algorithms in Java, 6th Edition,
2014.

- **Chapter 10. Hash Tables, Maps, and Skip Lists**
- **Section 10.3 The Sorted Map ADT**
- **pp. 396-401**

- **Chapter 11. Search Tree Structures**
- **Sections 11.1-11.2**
- **pp. 423-442**

Learning Objectives


- ◆ To be able to understand and describe the Sorted Map ADT;
- ◆ To be able to analyze the complexity of the Sorted Map ADT methods;
- ◆ To be able to implement the Sorted Map ADT with a binary search tree;
- ◆ To be able to explain the update operations for a binary search tree;
- ◆ To be able to apply the Sorted Map ADT and binary search tree.

Motivation

Suppose you have an array of integers and want to search for a particular integer k , how will you do it?

Motivation

- ◆ If the array is not sorted, then we need to scan all elements, hence $O(n)$.
- ◆ If the array is sorted, then we can do better. How?



Ordered Maps

- ◆ Keys are assumed to come from a total order.
- ◆ Items are stored in order by their keys
- ◆ This allows us to support **nearest neighbor queries**:
 - ◆ Item with largest key less than or equal to k
 - ◆ Item with smallest key greater than or equal to k

The Sorted Map ADT (Sec.10.3)

The Sorted Map ADT includes all methods of the Map ADT, plus the following.

- ❑ **firstEntry**(): Returns the entry with smallest key value (or null, if the map is empty).
- ❑ **lastEntry**(): Returns the entry with largest key value (or null, if the map is empty).
- ❑ **ceilingEntry**(k): Returns the entry with the least key value greater than or equal to k (or null, if no such entry exists).
- ❑ **floorEntry**(k): Returns the entry with the greatest key value less than or equal to k (or null, if no such entry exists).

The Sorted Map ADT


- ❑ **lowerEntry**(k): Returns the entry with the greatest key value strictly less than k (or null, if no such entry exists).
- ❑ **higherEntry**(k): Returns the entry with the least key value strictly greater than k (or null if no such entry exists).
- ❑ **subMap**(k_1, k_2): Returns an iteration of all entries with key greater than or equal to k_1 , but strictly less than k_2 .

Sorted Search Tables


- ◆ We store the map's entries in an array list A so that they are in increasing order of their keys.
- ◆ We refer to this implementation as a ***sorted search table***.

0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
2	4	5	7	8	9	12	14	17	19	22	25	27	28	33	37


Binary Search



- ◆ Binary search can perform nearest neighbor queries on an ordered map that is implemented with an array, sorted by key
 - similar to the high-low children's game
 - at each step, the number of candidate items is halved
 - terminates after $O(\log n)$ steps
- ◆ Example: **find**(7)



Sorted Search Tables



- ◆ A search table is an ordered map implemented by means of a sorted sequence
 - We store the items in an array-based sequence, sorted by key
 - We use an external comparator for the keys
- ◆ Performance:
 - Searches take $O(\log n)$ time, using binary search
 - Inserting a new item takes $O(n)$ time, since in the worst case we have to shift n items to make room for the new item
 - Removing an item takes $O(n)$ time, since in the worst case we have to shift n items to compact the items after the removal
- ◆ The lookup table is effective only for ordered maps of *small size* or for maps on which *searches* are the most common operations, while insertions and removals are rarely performed (e.g., credit card authorizations)

Performance of a sorted map

(implemented using a sorted search table)

Method	Running Time
size	$O(1)$
get	$O(\log n)$
put	$O(n)$; $O(\log n)$ if map has entry with given key
remove	$O(n)$
firstEntry, lastEntry	$O(1)$
ceilingEntry, floorEntry, lowerEntry, higherEntry	$O(\log n)$
subMap	$O(s + \log n)$ where s items are reported
entrySet, keySet, values	$O(n)$


Motivation

- ◆ Binary search on ordered arrays is efficient: $O(\log_2 n)$
- ◆ However, insertion or removal of an item in an ordered array is slow: $O(n)$
- ◆ Ordered arrays are best suited for *static* searching, where search space does not change.
- ◆ Binary search trees can be used for *efficient dynamic searching*.

Binary Search Trees

A binary search tree is a *proper* binary tree storing keys (or key-value entries) at its *internal nodes* and satisfying the following property:


- Let u , v , and w be three nodes such that u is in the left *subtree* of v and w is in the right *subtree* of v . We have $key(u) \leq key(v) \leq key(w)$
- Assuming there are no duplicate keys, we have $key(u) < key(v) < key(w)$




Binary Search Trees

External nodes do not store items


- and likely are not actually implemented, but are just null links from the parent




Is this a binary search tree?



Is this a binary search tree?




Is this a binary search tree?



Search

- ◆ To search for a key k , we trace a downward path starting at the root
- ◆ The next node visited depends on the comparison of k with the key of the current node
- ◆ If we reach a leaf, the key is not found
- ◆ Exercise: Write down the pseudo-code of the Search method.

Algorithm *Node TreeSearch(Key k , Node n)*



Search

Algorithm *Node TreeSearch*(Key k , Node n)

if $n.isExternal()$ // or “if $n == null$ ”

return *null*

if $k < n.key()$

return *TreeSearch*(k , $n.left()$)


else if $k = n.key()$

return n

else // $k > n.key()$


return *TreeSearch*(k , $n.right()$)

- ◆ Example: *get*(4):
 - Call *TreeSearch*(4, root)
- ◆ The algorithms for nearest neighbor queries are similar




Fundamental Property of Binary Search Trees

- ◆ What is an inorder traversal of a tree?
- ◆ Exercise: what does an inorder traversal of the following search tree produce?




Fundamental Property of Binary Search Trees

An inorder traversal of a binary search tree visits the keys in increasing order.



Fundamental Property of Binary Search Trees

- ◆ How to access the minimal key?
- ◆ How to access the maximal key?
- ◆ How to access the largest key less than or equal to k ?
- ◆ How to access the smallest key greater than or equal to k ?



Insertion


- ◆ Exercise: How to do insertion, *put*(k, o)?
- ◆ Have to insert k where a *get*(k) would find it.
- ◆ So natural that *put*(k, o) starts with *get*(k)

Insertion



- ◆ Exercise: How to do insertion, *put*(k, o)?
- ◆ We search for key k (using TreeSearch)
- ◆ Assume k is already in the tree then just replace the value.
- ◆ Otherwise, let w be the leaf reached by the search, we insert k at node w and expand w into an internal node

Insertion

◆ Example: insert 5




Insertion



Deletion



◆ How can we perform the operation *remove(k)*?

- E.g., remove 7, 5, 4, or 2



Deletion



- ◆ To perform operation *remove*(k), we search for key k
- ◆ Assume key k is in the tree, and let v be the node storing k
- ◆ If node v has a leaf child w , we remove v and w from the tree with operation *removeExternal*(w), which removes w and its parent
- ◆ Example: remove 4



Deletion (cont.)

- ◆ We consider the case where the key k to be removed is stored at a node v whose children are both internal
 - we find the internal node w that follows v in an **inorder traversal**
 - we copy $key(w)$ into node v
 - we remove node w and its left child z (which must be a leaf) by means of operation *removeExternal*(z)

◆ Example: remove 3



Self-study: can we let w be the internal node that precedes v in an inorder traversal? Then how to perform "delete"? pp. 428-429

Performance

- ◆ Consider an ordered map with n items implemented by means of a binary search tree of height h
 - the space used is $O(n)$
 - methods **get**, **put** and **remove** take $O(h)$ time
- ◆ The height h is $O(n)$ in the worst case and $O(\log n)$ in the best case

